
Psychometrics with AI Foundation Models

Haoran Ye¹, Tianrui Huo², Zheng Guo^{3,4},
Xin Zhang^{3,4}, Haonan Dong¹, Guojie Song^{1,✉}

¹State Key Laboratory of General Artificial Intelligence,
School of Intelligence Science and Technology, Peking University

²Yuanpei College, Peking University

³School of Psychological and Cognitive Sciences, Peking University

⁴Key Laboratory of Machine Perception (Ministry of Education), Peking University

hrye@stu.pku.edu.cn gjsong@pku.edu.cn

Abstract

Psychometrics is a subfield of psychology concerned with the objective measurement of latent constructs, such as values, personality, and intelligence. AI Foundation Models (FMs), notably large language and multimodal models, are reshaping this field by capturing rich linguistic and behavioural regularities at scale. This paper reviews the emerging synergy between FMs and psychometrics. It maps practical applications of FMs across the measurement pipeline, describes key methodologies for enhancing FM performance in psychometric contexts, and examines the theoretical implications of FMs for this discipline. In addition, we chart risks and offer actionable recommendations for the effective, rigorous, and ethical implementation of FMs in psychometric research and practice. Realizing the potential of FMs in psychometrics demands interdisciplinary collaboration and continued alignment between advances in computational and measurement sciences.

CONTENTS

1	Introduction	3
2	Applications of Foundation Models in psychometrics	3
2.1	Construct modeling and instrument development	3
2.1.1	<i>Semantic embeddings for estimating latent structures</i>	4
2.1.2	<i>Automatic item and stimulus generation</i>	5
2.1.3	<i>Pre-estimating item parameters</i>	7
2.1.4	<i>Translation and cross-cultural adaptation</i>	7
2.2	LLM-based response simulation	7
2.3	Estimating psychological states and traits	9
2.4	Adaptive and interactive assessment	9
2.4.1	<i>Adaptive assessment</i>	9
2.4.2	<i>Interactive assessment</i>	10
2.5	Score interpretation and intervention	10

3	Enhancing Foundation Models for psychometrics	10
3.1	Prompt engineering	10
3.2	Fine-tuning	12
3.3	Agentic AI	14
4	Theoretical implications of Foundation Models for psychometrics	15
4.1	Reshaping psychological taxonomies	15
4.2	Discovering and re-evaluating latent structures	15
4.3	Modeling cognition and representing populations	16
4.4	LLM psychometrics	16
5	Recommendations	17
5.1	Conceptual and methodological rigor	17
5.2	Robustness and reliability	18
5.3	Responsible and ethical integration	18
6	Summary and future directions	19
A	Terminologies	21

1 Introduction

Psychometrics is a scientific discipline concerned with quantifying latent psychological constructs such as values, personality, and cognitive abilities [85, 115]. It widely supports crucial applications across education, clinical assessment, organizational decision-making, social policy, and beyond.

Psychometrics usually involves mapping responses from instruments (e.g., self-report questionnaires, structured interviews, and performance tasks) onto statistical models, and yields interpretable estimates of psychological constructs [15]. Instrument development has traditionally relied on iterative cycles of item drafting, expert review, and empirical testing to establish reliability and validity [96]. This principled workflow is costly and slow, and it struggles to keep pace with research and applications that are expanding in scope, complexity, and diversity [33]. Computational approaches have offered new leverage that promise gains in efficiency, precision, and scalability [146]. For instance, supervised natural language processing (NLP) models enable automated scoring of digital traces [64]; topic models and distributional semantics facilitate content analysis [144]; and knowledge tracing and cognitive diagnosis support adaptive testing [8].

Earlier computational and machine learning (ML) models used in psychometrics are mostly discriminative and task-specific. They are trained on annotated datasets to classify or predict outcomes within narrow domains. A profound shift has occurred with the advent of *AI foundation models (FMs)*, such as *large language models (LLMs)* [1, 7, 14, 44, 138]. They differ remarkably from prior computational approaches in their deep semantic understanding, strong generative capabilities, and generality. Whereas earlier models typically rely on shallow features such as bag-of-words, FMs encode rich, contextualized semantic representations. They allow the models to better capture latent psychological constructs and handle ambiguity with human-like interpretation. Second, beyond purely discriminative tasks, many FMs can generate multimodal and human-like content. This enables new possibilities such as respondent simulation. Third, FMs are versatile across diverse tasks via instruction following and in-context learning. They can readily perform new tasks simply by following user instructions, without additional training.

Therefore, FMs promise to systematically transform the psychometric pipeline. Recent advances enable more accurate content analysis [106], introduce novel forms of automation [158], provide generalist assessment tools [160], support innovative measurement formats [155], and position AI itself as a subject of measurement [157]. At the same time, this shift introduces new challenges. For instance, additional sources of variance (prompting, decoding strategies, and model version) now interact with human variability [9, 118]. As a result, new controls and new forms of validity evidence (stability under paraphrase, cross-model agreement, and alignment with human data) are needed [157, 172]. The expanded capabilities and broader applications of FMs also warrant careful consideration of their limitations, as well as their ethical and social implications [18, 147].

This review charts the emerging integration of FMs into psychometrics (Fig. 1). We begin by surveying the applications of FMs to psychometrics (§ 2). Then, we examine methods for adapting and enhancing these models for psychometric tasks (§ 3). Next, we explore the theoretical implications that FMs pose for psychometrics (§ 4). To support responsible research and deployment, we offer recommendations for rigorous evaluation and ethical use (§ 5). Throughout, an interdisciplinary perspective is adopted to bridge advances in computational methods and measurement science, ensuring broad accessibility for a cross-disciplinary audience. Key terminologies are summarized in Appendix A.

2 Applications of Foundation Models in psychometrics

FMs are increasingly being integrated into psychometric pipelines. This section reviews these applications and organizes them according to the key stages of the measurement pipeline. We illustrate these applications in Fig. 2 and provide a detailed summary with representative examples in Table 1.

2.1 Construct modeling and instrument development

Psychometric instrument development often begins with rigorous construct conceptualization. Then, it proceeds through operationalization, item generation, and validation. FMs are now being used to




Figure 1: The structure of this review.

support this process holistically. Semantic embeddings help map and diagnose the construct space and reveal taxonomic issues such as jingle-jangle fallacies (§ 2.1.1). LLMs assist blueprint-constrained item and stimulus generation (§ 2.1.2). Models can pre-estimate item parameters to triage large item pools prior to piloting (§ 2.1.3). Also, AI systems aid translation and cross-cultural adaptation (§ 2.1.4).

2.1.1 Semantic embeddings for estimating latent structures

Semantic embeddings are numerical vectors that represent multimodal content in a high-dimensional space, where geometric proximity approximates semantic relatedness. These modern representations capture nuanced, context-dependent meaning, and can be compared using simple metrics such as cosine similarity.

Recent work shows that embedding similarity between test items can meaningfully predict their psychometric structure. For example, it predicts a priori factorial structures of personality tests with moderate-to-high correlation to human-derived structures [97]. For some scales, embedding similarity is shown to correlate with empirically observed item-item correlations ($r = 0.67$), and can help flag potential sources of model misfit [32] or abbreviate scales [63]. Principal component analyses (PCA) of text embeddings are shown to explain more cumulative variance than PCA of human responses [97]; some studies report near-perfect consistency between factor loading matrices derived from embeddings and from survey ratings [26]. These findings echo pre-LLM evidence that semantic relationships can account for stable portions of variance and, in some instruments, large fractions of observed item-item correlations [5].

LLM embeddings also enable extrapolation and integration. Research scales up analyses of personality lexicals with embeddings and provides evidence supporting two- or three-factor models of personality




Figure 2: Overview of foundation model applications in psychometrics.

[26], in contrast to the more prominent big five model [38]. Interestingly, applying singular value decomposition (SVD) to LLM log-probabilities of trait-descriptive adjectives has also been reported to rediscover the big five traits [132]. These discrepancies are likely due to the sensitivity of LLMs to implementation choices [9]. At a broader level, mapping items, scales, and construct labels into a shared semantic space offers a scalable method for diagnosing potential taxonomic problems, such as the "jingle-jangle" fallacies that have long challenged the field [52, 151]. The theoretical implications of this capability are discussed further in § 4.

Semantic embeddings extend beyond text to encompass multimodal content. AI researchers utilize multimodal embeddings [113] to quantify the similarity between images and textual construct labels, which allows them to evaluate the value orientations of multimodal FMs [76].

Overall, embeddings provide powerful pre-data diagnostics and serve as a valuable scaffold for construct modeling and scale design. However, they should be viewed as a complement—not a substitute—for empirical validation and theoretically grounded interpretation. For example, research highlights the potential of human-AI collaboration in assessing item content validity [98]. It is important to note that semantic similarity may capture surface features such as wording or stylistic choices rather than underlying latent constructs, which can artificially inflate inter-factor correlations [31, 72]. Such deficiencies become apparent when embeddings fail to yield consistent results for reverse-scored items [63].

2.1.2 Automatic item and stimulus generation

The development of psychometric items and stimuli—ranging from test questions to evocative images—has long relied on labor-intensive processes and domain expertise. Automatic item generation (AIG) promises to alleviate this burden [36]. FMs have rapidly shifted AIG from hand-crafted templates to scalable, psychometrics-aware, and blueprint-constrained drafting with minimal human intervention.

Initial applications in language assessment demonstrate FMs capable of generating complete reading-comprehension test batteries [6]. This approach has since expanded to cognitive reasoning tasks, including natural language inference. Strategically prompted LLMs produce items with psychometric properties comparable to expert-authored baselines [71]. The scope has broadened beyond multiple-choice formats to encompass situational judgment tests and open-ended stimuli for personality and

Table 1: Overview of how FMs are used in psychometrics.

Application Area	Description & Key Benefits	Representative Examples
Semantic embeddings for internal structure & diagnostics (§ 2.1.1)	Numerical representations capture semantic relationships between items, predict factorial structures, and enable cross-measure comparisons.	Embedding similarity predicts item-item correlations and factorial structures [97]; Mapping constructs into shared semantic space reveals jingle-jangle fallacies and addresses taxonomic incommensurability [151].
Automatic item & stimulus generation (§ 2.1.2)	Scalable creation of psychometrically sound test items, from multiple-choice questions to visual stimuli.	Creative and personality assessment items with retained psychometric properties [60, 73]; Emotionally evocative images rated as indistinguishable from authentic photographs [86].
Pre-estimating item parameters (§ 2.1.3)	Predict item difficulty, discrimination, and response time prior to human administration, streamlining test development.	Direct LLM prediction of IRT difficulty parameters [55]; LLM ensembles as synthetic respondents for IRT-based analyses [84].
Translation & cross-cultural adaptation (§ 2.1.4)	Automated translation and cultural adaptation of psychological instruments while preserving construct validity.	Accurate, conceptually equivalent cross-cultural adaptations [128]; LLM translations outperforming traditional methods and approaching expert-level quality [66].
LLM-based response simulation (§ 2.2)	LLMs prompted to respond as specific individuals or groups across psychological measures, enabling rapid prototyping and rare scenario exploration.	Persona-prompted LLMs maintaining behavioral consistency across contexts [2]; GPT-4 predicting 91% of treatment effect variation across 70 experiments [3].
Estimating psychological states and traits (§ 2.3)	Measure psychological constructs from unstructured data (social media, speech, video) using multimodal FMs.	Personality estimation from social media posts [106]; Multimodal trait estimation integrating dialogue, speech, and video [75].
Adaptive assessment (§ 2.4.1)	Dynamic item selection and cognitive diagnosis that adapts to individual responses in real-time, improving measurement efficiency.	TestAgent analyzing open-ended responses for adaptive testing [160]; Multi-agent systems for optimal question selection [21].
Interactive assessment (§ 2.4.2)	Transform static assessments into engaging dialogues or game-like experiences while maintaining psychometric rigor.	PsyDI conducts psychological profiling through free-form conversations [78]; PsychoGAT embeds validated scales within interactive fiction games [155].
Score interpretation & intervention (§ 2.5)	Agentic systems that integrate assessment findings and generate personalized recommendations for treatment or intervention.	Psy-LLM conducts guided interviews and generates personalized suggestions [68]; PsychBench challenges models to prioritize medications and recommend treatment plans [82].

creative assessment, and the item qualities are shown to be retained [60, 73]. Large-scale evaluations in higher education demonstrate that LLM-generated items can achieve psychometric performance comparable to expert-created questions for standardized examinations [54]. Further advancement introduces a comprehensive framework leveraging network psychometrics to generate items that are simultaneously diverse, theoretically grounded, and structurally robust [119].

Beyond text-based items, FMs enable the generation of emotionally evocative visual stimuli for psychological research. Orchestrating LLMs and vision FMs, researchers can create images that are rated as nearly indistinguishable from authentic photographs, demonstrating the potential of high-fidelity synthetic visual stimuli [86].

This progress extends to the evaluation of AI itself. As LLMs exhibit increasingly human-like characteristics, researchers have begun evaluating the psychological traits of these models using items that are pre-authored or dynamically generated by LLMs [23, 117, 158]. Parallel efforts focus on training LLMs to serve as sophisticated item generators, capable of producing novel and discriminating items that adapt to or extrapolate from specified difficulty parameters [57, 141].

2.1.3 Pre-estimating item parameters

Pre-estimating the psychometric properties of test items—such as Item Response Theory (IRT) difficulty, discrimination, and response time—prior to human administration can streamline test development and reduce piloting demands. LLMs facilitate this process and enable efficient, low-cost screening of large pools of candidate items.

Two primary strategies have emerged for leveraging LLMs in this context. The first involves using LLMs to directly predict psychometric properties [55] or to extract linguistic features that inform subsequent ML models [116]. The second approach treats LLMs as synthetic respondents, generating simulated response data for IRT-based analyses [135]. However, without additional training or calibration, direct LLM-based scoring is often unstable and misaligned with human data [55, 135]. Ensembles of LLMs [84] or models specifically trained for IRT alignment [121] can better approximate item parameters.

2.1.4 Translation and cross-cultural adaptation

Adapting psychometric instruments across cultures and languages requires preserving validity and equivalence. Recent advancements include AI-assisted translation, where systems evaluate and refine human-generated translations [45], as well as fully automated translation by LLMs [124, 128]. Notably, expert-blind evaluations indicate that LLM-based translations can outperform traditional methods and approach expert-level quality [66]. Nevertheless, translation performance declines substantially for low-resource languages and is influenced by the genetic similarity between source and target languages [66]. Approaches like multicultural data augmentation and culture-aware fine-tuning offer promising solutions to these challenges [74].

2.2 LLM-based response simulation

Psychometric data collection traditionally relies on recruited participants and repeated field trials—a process that is slow, expensive, and limited in scope. Emerging research therefore explores using LLMs to simulate human respondents across diverse psychological contexts. This approach promises a scalable complement to human data collection, enabling researchers to probe rare scenarios, explore counterfactual design choices, and rapidly iterate on instrument development. The premise is straightforward: LLMs can be prompted with contextual cues and instructed to respond as specific individuals would across psychological measures. However, the fundamental question remains whether synthetic responses can meaningfully substitute for human psychological data.

Current evidence presents a mixed picture. On one hand, models demonstrate remarkable capabilities in maintaining behavioral consistency across contexts. Persona-prompted LLMs can produce responses that align with established personality structures, sustain stable trait-like patterns across repeated assessments, and exhibit theoretically coherent relationships between different psychological constructs [2, 83, 122]. When scaled to population levels, synthetic respondents can approximate subgroup-consistent response patterns and recover broad conceptual landscapes observed in human samples, particularly when conditioning cues are carefully calibrated to empirical priors [4]. Recent

large-scale evaluations provide encouraging evidence [3]: GPT-4 predicts approximately 91% of the variation in average treatment effects across 70 preregistered, U.S.-representative social-science experiments [48]; fine-tuned LLMs can outperform existing cognitive models in predicting individual decision making [12]; and large-scale agent simulations (N=1,052) built from interview transcripts successfully predict participants' survey responses with 85% of the accuracy [103].

However, research reveals systematic limitations that challenge the psychometric utility of this methodology [3]. First, *diversity*: LLM outputs are often too generic or stereotyped, failing to capture the rich variation observed in human populations. For example, LLMs produce homogeneous and narrowly clustered responses in the 11-20 money request game, while human participants show a more diverse response distribution [35]. Second, *demographic bias*: models exhibit systematic inaccuracies for certain demographic groups, harmfully misportraying and flattening the representations of identity groups [147]. Third, *sycophancy*: models attempt to please users or provide socially desirable responses [120, 130]. For example, LLMs tend to respond affirmatively to both original and reverse-coded items [136]. Fourth, *alienness*: results may appear human-like on the surface but arise from fundamentally non-human mechanisms, leading to empirical artifacts such as over-correlated personality dimensions [108]. Fifth, *generalization*: accuracy deteriorates in out-of-distribution settings involving novel measures, instruments, or populations.

These limitations are compounded by the brittleness of simulated responses to seemingly minor variations in prompting or model configuration. Studies document substantial shifts not only in response means but also in correlation structures and psychometric properties when prompts are reworded or when different model versions are employed [13, 125]. This sensitivity undermines the stability that is fundamental to psychological measurement and raises questions about the replicability of findings based on synthetic data. LLMs also exhibit concerning response sets. When responding to multiple-choice questions with minimalist prompts, simulations display systematic ordering biases (e.g., a preference for the first option) and labeling biases (e.g., a preference for option "A") that, after adjustment, result in uniformly random option distributions [29].

The fidelity and utility of response simulations appear to depend on several methodological and experimental factors. *Prompting strategies* play a crucial role: specifying clear, detailed personas with rich contextual information (e.g., backstories and values) substantially improves simulation accuracy compared to simple or abstract prompts [4, 51, 80]. Context-rich approaches, such as incorporating hour-long interview transcripts, can better anchor simulated individuals [103]. Evidence suggests that framing the task as expert forecasting, rather than role-playing, is beneficial [48]; instead of prompting, "You are a person with X characteristics," researchers can ask, "What proportion of people with X characteristics would respond Y?" This "LLM-as-expert" framing can help counter sycophancy and social desirability bias [3].

Test format also influences simulation quality. LLMs perform better on structured, judgment-based tests with clear cues (e.g., multiple-choice questions, scenario ratings) [83] and should be avoided for unstructured or projective tests that require deep, embodied understanding or creative expression [35]. *Model selection* matters as well: base models that have not undergone extensive safety alignment better capture the full distribution of human responses, as safety alignment trains models to produce socially desirable and normative responses, suppressing heterogeneity and inducing bias [3, 22, 120]. Fine-tuning models on curated social-science datasets can mitigate assistant-style biases and improve simulation accuracy [3, 80]. Finally, the *underlying task* shapes appropriateness: simulations are most suitable for early-stage prototyping (e.g., generating item candidates, testing questionnaire wording, exploring theoretical boundaries) and non-interactive, judgment-only uses [28], while strategic or interactive tasks and socially sensitive topics pose greater challenges [35]. Direct simulation of minority or marginalized groups as substitutes for real human data raises significant ethical concerns and validity issues [13].

Looking ahead, LLM-based social simulations are already useful for pilot studies, exploratory research, and sensitivity analyses to guide study design before collecting human data [3]. As accuracy improves, these simulations may enable exact replications of prior studies and, eventually, large-scale or counterfactual experiments that would be impractical or impossible with human participants. Realizing this potential requires ongoing validation, standardized evaluation protocols, and careful mapping of where simulations succeed or fail. For a more comprehensive discussion, see the referenced position paper [3].

2.3 Estimating psychological states and traits

Psychometric scoring transforms responses to measurement items into numerical indicators of underlying latent constructs. Traditionally, scoring is based on predetermined keys applied to responses from standardized and validated instruments, where each item has an established relationship to the construct. With digital-behavioral data, the functional equivalent of "items" may be traces from everyday activity—such as language use, speech patterns, or video-based behavioral cues—that can then be linked to constructs through ML models. FMs extend this paradigm by leveraging powerful multimodal representations and mapping complex, unstructured inputs onto quantitative construct estimates.

Broadly, three methodological families have emerged. First, LLM-as-a-judge approaches prompt LLMs to read text (for example, social media posts, human-AI conversations, and interviews) and output state or trait estimates [106, 107, 170]. Zero-shot estimation yields modest but meaningful convergent validity with self-reports on personality tests (typical $r \approx 0.2-0.3$), which improves when interactions are designed to elicit diagnostic content (mean $r \approx 0.44$ in elicitation-focused chats) [107]. The methods can benefit from few-shot prompting (incorporating a few solution examples in prompts for LLMs to learn in situ), reasoning-enhancing prompting [104, 123], and fine-tuning on labeled datasets [152].

Second, representation-based approaches leverage FM embeddings as features for downstream ML-based classification or regression. This hybrid strategy yields lower prediction error for personality traits [90] and mental health [114] compared to zero-shot LLM inference. Notably, supplementing these embeddings with additional linguistic features does not further enhance performance, suggesting that FM representations inherently capture key linguistic characteristics (e.g., those compiled by LIWC [16]) without the need for explicit feature engineering.

A third line of research seeks to align the reasoning processes and outputs of LLMs with traditional psychometric measures, thereby strengthening both theoretical foundations and empirical validity. In one approach, LLMs are prompted to rate established items based on the content to be assessed (e.g., social media posts), and these ratings are subsequently aggregated to yield construct-level scores [127, 154, 156]. This procedure enables the calculation of internal consistency in a manner directly analogous to conventional psychometric assessments. Alternatively, LLMs can be employed to extract novel item-like statements (e.g., generating statements such as "I love parties" from posts expressing enthusiasm for social gatherings) [158]. The resulting measurements are then validated against external criteria, such as established theoretical frameworks, trait structures, and prior computational methods (e.g., theory-driven dictionaries).

The applications of FMs in state/trait estimation extend beyond textual content. Multimodal FMs integrate information from multiple cues, such as dialogue transcripts, speech, and video, to improve prediction [75, 167]. After task-specific fine-tuning, FMs are also able to capture subtle cues from facial expressions, enabling emotion recognition [165] and the diagnosis of depression [59].

2.4 Adaptive and interactive assessment

Traditional psychological assessment relies on standardized instruments that present the same questions to all individuals, regardless of their unique characteristics or response patterns. This one-size-fits-all approach can be inefficient; it potentially over-tests some individuals while under-sampling others, and it often fails to engage participants meaningfully. FMs are now enabling a fundamental shift toward dynamic, personalized assessment systems that adapt in real-time to individual responses and contextual factors.

2.4.1 Adaptive assessment

Adaptive assessment tailors the assessment experience to each individual. These systems operate through three interconnected components that work together to optimize measurement efficiency and accuracy. First, cognitive diagnosis algorithms continuously estimate the target psychological construct based on accumulated responses, updating their understanding as new information becomes available. Second, question selection algorithms identify the most informative items to present next, based on the current estimates. Third, comprehensive item banks provide the reservoir of questions from which selections are made.

FMs promise to enhance each of these components. For cognitive diagnosis, LLM agents can now analyze responses to open-ended questions, detecting nuanced patterns and anomalies that were previously inaccessible to automated systems [160]. This capability extends adaptive assessment beyond traditional multiple-choice formats to include rich, naturalistic responses. For question selection, multi-agent systems can summarize response histories, reason about optimal next questions, and dynamically refine their selection strategies [21]. The question banks themselves are expanding through automated item generation, which can scale the creation of passages, questions, and conversational scenarios to unprecedented levels (see § 2.1.2). These advances collectively enable more efficient and personalized assessments.

A parallel development involves using adaptive assessment principles to evaluate LLMs themselves. Often under IRT frameworks, researchers dynamically calibrate item difficulty based on AI performance and weighting items by their inferred difficulty, aiming to achieve accurate evaluations with smaller test sizes and more discriminative items [43, 69, 110].

2.4.2 Interactive assessment

Beyond adaptation, FMs enable entirely new forms of interactive assessment that transform the assessment experience from a passive question-answering session into an engaging dialogue. This shift addresses longstanding critiques of traditional assessment methods, which are often characterized by their static nature, lack of personalization, and limited participant engagement. Moreover, it democratizes access to sophisticated psychological evaluation, making assessment tools available beyond traditional clinical or research settings.

Current approaches to interactive assessment fall along a spectrum of structure and openness. At one end, frameworks transform established psychological scales into personalized, engaging interactive fiction games. This approach maintains the psychometric rigor of closed-ended assessment while dramatically enhancing user experience [155]. At the other end, researchers employ conversational AI systems that measure psychological constructs through free-form, multi-turn conversations. This approach allows for comprehensive profiling while maintaining the natural flow of human dialogue [78].

Two approaches balance comprehensiveness and nuance (free-form dialogue) against the higher reliability and validity (structured, game-like test formats). Encouragingly, early evidence suggests that, regardless of the specific approach, participants consistently report higher engagement and satisfaction with interactive assessment systems compared to conventional tests [155].

2.5 Score interpretation and intervention

The value of psychometric assessment extends beyond generating numerical scores to encompass the interpretation of results and recommendation of actionable interventions. Emerging attempts orchestrate FMs to build agentic systems that can conduct guided interviews, integrate findings across items and scales, and generate personalized suggestions for symptom management or referral [50, 68]. Recent benchmarking tasks further challenge models to prioritize medications or recommend treatment plans based on detailed clinical vignettes, underscoring growing expectations that AI should facilitate not only diagnosis but also care planning [82]. However, these approaches remain in their infancy, with limited validation of their clinical utility and significant ethical concerns about AI systems' making treatment recommendations without human oversight.

3 Enhancing Foundation Models for psychometrics

FMs demonstrate strong capabilities out of the box, but their performance can be further improved through enhancement strategies. This section presents methods for optimizing FMs in psychometric contexts.

3.1 Prompt engineering

FM performance is inherently dependent on the input context [94]. Thoughtful prompt engineering can unlock their superior psychometric utility. Broadly, prompts in psychometrics serve as controllers

Table 2: Overview of enhancement strategies for Foundation Models in psychometrics.

Enhancement Strategy	Technique & Key Benefits	Representative Examples
Prompt engineering (§ 3.1)	<i>Prompts as controllers:</i> Enforce constraints on FM outputs (e.g., item attributes, construct definitions, scoring rubrics).	Specifying item attributes (e.g., topic, difficulty) for item generation [6]; Incorporating analytical rubrics to improve scoring accuracy [150].
	<i>Prompts as scaffolds:</i> Provide structured guidance or intermediate steps to improve model performance.	Guiding models with psychological questionnaires as a chain-of-thought (PsyCoT) to improve personality detection from text [156].
	<i>Prompts as levers:</i> Actively manipulate and probe FM behavior for role-playing, improving simulation fidelity, and enhancing measurement reliability and validity.	Prompting LLMs to act as a "psychologist" and "game designer" to convert scales into interactive fiction games [155]; Using detailed, behaviorally specific descriptions to improve response simulation fidelity [56].
Fine-tuning (§ 3.2)	<i>For semantic embeddings:</i> Match cosine similarity of item embeddings with empirical item-item correlations.	Fine-tuning embedding models using IPIP dataset to make the cosine similarity between item embeddings match empirical item-item correlations [151].
	<i>For item generation & parameter estimation:</i> Control construct fidelity and difficulty, screen generated items, and pre-calibrate item parameters.	Fine-tuning on the IPIP dataset for better personality-specific item generation [49]; Training FM evaluators to screen generated items [40].
	<i>For response simulation:</i> Calibrate models at individual, group, or distributional levels to predict responses.	Training models with distributional calibration to output the full distribution of population responses, conditioned on group characteristics [20].
	<i>For measuring psychological properties:</i> Predict psychometric targets from free-form content (e.g., text, multimodal data) using end-to-end or hybrid approaches.	End-to-end fine-tuning on counseling dialogues to predict psychometric scores [154]; A hybrid pipeline using FM embeddings from social media to estimate personality [90].
	<i>For adaptive assessment:</i> Train parameter estimators and question generators using calibrated IRT parameters as supervision signals.	Using reinforcement learning to train an item generator that is rewarded for producing items matching target IRT parameters [141].
Agentic AI (§ 3.3)	<i>Single-agent systems:</i> Use tools, interact with the world, and adapt behavior based on feedback.	A single agent adapts to user answers by selecting items and updating profiles iteratively [78]; Using Retrieval-Augmented Generation (RAG), where the agent retrieves expert reasoning for personality detection [79].
	<i>Multi-agent systems:</i> Coordinate multiple specialized agents to collaboratively achieve complex psychometric outcomes.	"TestAgent" uses multiple agents for adaptive assessment, including selecting questions, eliciting answers, and generating diagnostic reports [160].
	<i>Multimodal orchestration:</i> Chain and integrate FMs across modalities (text, audio, visual) to enable more versatile capabilities.	A workflow combining language and visual FMs generates emotionally evocative images, with agents responsible for prompt formation, description diversification, and image synthesis [86].

(to enforce constraints), scaffolds (to elicit capabilities), and levers (to shape synthetic respondent behavior).

Prompts as controllers explicitly instruct and constrain FM outputs for specific psychometric purposes. In automatic item generation, practitioners can effectively manage item attributes by clearly specifying the topic, difficulty level, item format, and design features within their prompts [6, 11, 95]. When combined with carefully chosen, authoritative examples ("few-shot" learning), this approach allows for more precise control over both the difficulty and stylistic characteristics of generated items [6, 99]. For measurement tasks, practitioners achieve better results by providing comprehensive, operationalized definitions of the constructs they wish to assess, rather than relying solely on construct names [117, 143]. Similarly, incorporating high-quality analytical rubrics that mirror human grading logic enhances the scoring accuracy of LLMs [150].

Prompts as scaffolds provide structured guidance or intermediate steps to support and enhance the model's reasoning, ultimately yielding better psychometric outcomes. AI practitioners typically employ scaffolding prompts to plan steps (planning), think aloud (reasoning), check the model's own output (self-reflection), and iteratively improve generation (self-improvement). For example, a popular and generic approach is to instruct models to "think step-by-step," incorporating deliberative reasoning (e.g., chain-of-thought) that mirrors human System 2 processing [148]. Some psychometric applications also employ this generic strategy [79, 170], often resulting in notable improvements. Others utilize more domain-specific scaffolds [88, 156]. For example, researchers introduce using psychological questionnaires as chain-of-thought (PsyCoT), which guides models through a structured assessment process: first generating estimated scores for individual scale items, then aggregating these scores into final construct-level estimates [156]. A similar approach explicitly instructs the LLM to reason about the emotions and emotion regulation underlying texts before outputting the detected personality traits [79].

Prompts as levers actively manipulate and probe the behavior of FMs. In interactive assessment, LLMs are prompted to take on roles such as "psychologist" and "game designer," converting scales into interactive fiction. These role-playing prompts guide both content and tone, reportedly improving engagement while preserving reliability and validity [155].

In response simulation tasks, theory-informed and role-playing prompts can improve simulation fidelity [67, 122]. For example, researchers achieve this by expanding basic personality lexicons into detailed, behaviorally specific personality descriptions, enabling LLMs to adopt diverse yet theoretically consistent behaviors [56]. More comprehensive approaches incorporate empirically derived human belief networks, aligning role-playing LLMs across 64 topics and 9 latent factors [24]. In general, providing more detailed information about the target profile (e.g., hour-long interview transcripts) can maximize simulation fidelity [4, 30, 131]. Furthermore, the framing of the simulation task itself matters: "LLM-as-expert" prompts that ask models to forecast what people would do may yield more accurate predictions than "LLM-as-subject" role-playing prompts, as discussed in § 2.2.

Prompts as levers can also improve the reliability and validity of measurements. By employing a diverse set of prompt variants and aggregating their outputs, practitioners can more thoroughly assess and improve the reliability (e.g., internal consistency) of the resulting measurements [127]. Prompt variants also help address FM biases such as option order bias by shuffling the options randomly [29]. Distribution elicitation techniques, which request models to output full response distributions or proportions rather than single responses, offer another approach to counter sycophancy and improve the validity of population-level estimates [3].

3.2 Fine-tuning

FMs are pretrained on vast datasets, enabling them to perform a wide range of tasks out of the box. However, their performance on specialized tasks can often be further improved through fine-tuning, a process in which FMs are trained on smaller, task-specific datasets to better adapt to particular applications. Fine-tuning is commonly employed to enhance model performance on psychometric tasks.

When training semantic embeddings to capture psychometric relationships [32, 151], the objective is to ensure that the cosine similarity between item embeddings closely matches the empirical item-item correlations observed in human data. Typically, the embedding model takes psychometric items as input and outputs their high-dimensional representations. Its training maximizes the

correspondence between embedding similarities and ground-truth correlation matrices, often via regression-based losses. Open-Source Psychometrics Project [101] can be used to construct the training and validation dataset. Its data includes both the raw item content and the empirical correlation structure, facilitating supervised fine-tuning and robust benchmarking of the models.

In automatic item generation and parameter estimation, fine-tuning serves three main objectives. The first objective is *construct fidelity and difficulty control for item generation*. Fine-tuned models can follow detailed instructions more closely and can be adjusted to create items with specific levels of difficulty. For example, fine-tuning FMs on the International Personality Item Pool (IPIP) [39] enables better construct-specific item generation [47, 49]. Training both difficulty evaluators and question generators allows for better control over item difficulty and improves test efficiency [43, 69, 110, 142]. The second objective is *reliable and scalable screening of generated items*. Fine-tuned FM evaluators, trained to imitate human item evaluators, act as more reliable proxies to triage large candidate pools before human review [40, 99, 171]. The third objective is *accurate pre-calibration for parameter estimation*. Research fine-tunes FMs on historical responses (and their fitted IRT models) in hopes of simulating how test takers would have responded to unseen items [121, 162, 163]. As an alternative to response simulation, other approaches directly train FMs, or downstream ML predictors based on FM generated features, to predict IRT parameters from item content [19, 61, 77, 93, 116].

In response simulation, existing training approaches can be categorized by their level of calibration and optimization objective: individual-level calibration, group-level calibration by characteristics, and distributional calibration [58]. For *individual-level calibration*, researchers train models to capture and predict individual survey responses [58, 58] and behaviors in experiments [12]. These models are typically trained using standard supervised learning methods, taking as input individual behaviors or profiles and learning to predict unseen responses through imitation (e.g., maximizing the likelihood of the "true response"). For *characteristics-driven calibration*, the premise is that factors such as demographics and value preferences influence decision-making. Accordingly, models are trained to predict responses based on these characteristics [53, 111, 131]. However, this approach cannot capture intra-group variation. To address this, research on *distributional calibration* trains models to output the full distribution of population responses, with or without conditioning on group characteristics [20, 129]. These models are usually trained by minimizing the divergence between the predicted and ground-truth distributions.

Recent research advocates for individual-level modeling, as it provides more nuanced representations. In addition, while distributional data is highly demanding to obtain, it can be derived by marginalizing over group characteristics using individual-level models [131]. Taken together, these training approaches improve individual-level and group-level marginal alignment and enable some generalization across surveys, experiments, and groups. However, true generalization beyond the training data remains questionable and requires further validation [3, 25, 108, 157].

Training for measurement usually relies on datasets that pair free-form, ecologically valid content (e.g., counseling dialogues [154] and social media posts [158]) with either self-report psychometric scores or expert psychological annotations. Annotation granularity varies, ranging from item-level (e.g., associating brief textual descriptions with value valence or affective ratings) [158] to discourse-level (e.g., mapping entire blog posts to personality profiles or clinical categories) [90].

Some supervised fine-tuning approaches directly optimize FMs to predict psychometric targets from raw text, leveraging large-scale annotated corpora [154]. Alternatively, research has introduced two main hybrid pipelines. In one approach, FMs are first used to generate high-dimensional embeddings from free-form content, which are then fed into downstream ML models to estimate psychometric constructs [90]. In the other approach, theory-driven linguistic features are extracted from the free-form content, and their counts are provided as input to fine-tuned FMs for construct estimation [62]. These hybrid strategies often enhance computational efficiency and interpretability. In addition, the integration of multimodal and temporal data introduces additional technical complexity, necessitating the development of fusion architectures, specialized regression heads, and temporal modeling components [34, 75, 153].

Challenges persist despite methodological advances. Robustness to domain shift—generalization beyond the specific contexts represented in training data—remains limited. Systematic evaluation of reliability and validity at the level of established psychometric standards is rarely reported. Critically, the dependence on paired free-form content and gold-standard psychometric labels imposes

substantial costs of data collection and annotation, limiting scalability and the breadth of construct coverage.

Training for adaptive assessment typically integrates a parameter estimator and a question generator, both of which are also discussed individually above. The process generally begins by constructing psychometric targets (labels) from existing items. This involves calibrating items using response data from examinees (human, AI, or a hybrid of both) to estimate IRT parameters: primarily difficulty (b), and sometimes discrimination (a) and guessing (c), typically via expectation-maximization methods [27]. These calibrated parameters then serve as supervision signals for subsequent model training [139, 141]. As discussed in automatic parameter estimation, synthetic response patterns generated by FMs can augment or partially replace human data, thereby extending coverage and reducing costs [83, 89].

Using the constructed datasets, one can train a regression model to estimate item IRT parameters directly from item content. The model is typically optimized to minimize the mean squared error between predicted and empirical parameters, enabling accurate parameter estimation prior to field validation [141]. Building on this, a parameter-conditioned item generator can be trained to produce items that align with specified psychometric targets (e.g., difficulty). Training paradigms include both supervised learning, which encourages the generator to emulate curated item exemplars [89], and reinforcement learning, which rewards the generator for producing items whose estimated parameters (as predicted by the parameter prediction model) closely match the intended targets [141]. These two models, i.e., the parameter predictor and the item generator, can be trained sequentially or in a joint, co-evolutionary framework. The latter facilitates mutual refinement and improved alignment between psychometric calibration and item generation [57].

These training techniques are complementary and can be combined into multi-stage pipelines for advanced applications such as interactive, adaptive assessment and real-time psychometric assessment. For example, a system may integrate fine-tuned item generators, parameter predictors, and scoring models within an agentic framework to iteratively select, generate, calibrate items, and evaluate open-ended responses.

3.3 Agentic AI

Agentic AI refers to the design and deployment of multi-step, autonomous, and goal-directed AI systems capable of planning, reasoning, and collaborating to accomplish complex psychometric tasks. State-of-the-art agentic AI systems have shown promise in addressing complex scientific challenges, such as designing novel nanobodies [134], creating molecules with specific properties [161], and advancing mathematical frontiers [100]. In psychometrics, agentic AI is still in its early stages of development and is broadly explored in three ways: single-agent, multi-agent, and multimodal approaches.

Single-agent systems can use tools, interact with the external world, and iteratively adjust their behavior based on feedback. In psychometric contexts, these agents adapt to respondents' answers, selecting items, conversing with users, and updating profiles through ongoing, iterative processes [78, 145]. Another key advantage of agentic AI is its ability to retrieve knowledge from external sources. FMs encode static knowledge in their parameters and may not meet the needs of specialized tasks, especially those requiring domain expertise. In such cases, retrieving information from external sources is essential. Retrieval-augmented generation (RAG) in agentic AI injects knowledge by retrieving expert information from external databases using semantic matching. For example, one system retrieves expert-curated reasoning processes to guide LLMs in more accurate personality detection [79]. Similarly, systems for educational question generation, improvement, and evaluation ground their outputs in domain-specific course materials retrieved from relevant sources [99].

Multi-agent systems coordinate multiple AI agents, each assigned a distinct role, to collaboratively achieve psychometric outcomes that surpass the capabilities of single models. These systems break down complex measurement tasks into specialized subtasks, with each agent contributing complementary expertise. "Generative psychometrics" illustrates this approach through a systematic three-stage pipeline for value measurement [158]. First, specialized agents extract item-like observations from free-form text. Next, additional agents measure the values embedded in each observation. Finally, an aggregation step synthesizes these individual measurements into comprehensive construct scores. TestAgent, as another example, positions LLMs as adaptive assessment experts that select

personalized questions, elicit open-ended answers, generate similar or variant questions to address label or subjective uncertainty, detect anomalies, and produce diagnostic reports [160]. Multi-agent frameworks can also efficiently build and enrich domain knowledge graphs, enabling applications such as automated hypothesis generation and scientific discovery [87, 140].

Multimodal orchestration represents another line of research that leverages the capabilities of FMs across different modalities. Research on generating emotional images develops workflows that chain together language and visual FMs in coordinated stages: first, seed prompt formation agents generate descriptions of emotional images; next, prompt set generation agents create diverse descriptions targeting specific emotions; finally, text-to-image synthesis agents produce the resulting emotional stimuli [86]. Expert review is strategically interleaved throughout this pipeline to ensure both scalability and validity. Moreover, multimodal orchestration across FMs enables the fusion of text, audio, and visual information for trait measurement in richer, more contextually grounded settings [75, 164, 166].

4 Theoretical implications of Foundation Models for psychometrics

The integration of FMs into psychometrics extends beyond practical and methodological innovation; it also opens new frontiers for psychological theories.

4.1 Reshaping psychological taxonomies

The ability of FMs to map psychological constructs into a shared semantic space, as shown in § 2.1.1, goes beyond methodological convenience; it offers a data-driven approach to the long-standing challenge of taxonomic incommensurability in psychology [70, 149]. This approach can help uncover common taxonomic issues in the field, such as the jingle fallacy, where measures with similar labels actually assess different constructs, and the jangle fallacy, where different labels are used for measures that assess the same construct. By providing a scalable way to quantify relationships between measures, FMs enable researchers to systematically identify and address these problems.

The potential theoretical implication is a paradigm shift from fragmented, redundant, and top-down taxonomies toward a more unified, parsimonious, and empirically grounded psychological taxonomy. Recent research supports the use of trained embeddings to clarify conceptual confusion in psychology, suggesting that a smaller set of constructs may be sufficient to capture meaningful distinctions in Psychology [151].

4.2 Discovering and re-evaluating latent structures

FMs can serve as powerful engines for discovering and re-evaluating the latent psychological structures. Embedding-based analyses have not only rediscovered established structures such as the Big Five personality traits, but have also provided evidence for alternative models, including two- or three-factor models for extensive personality lexicons [26]. Factor structures derived from everyday human language use have been shown to predict non-questionnaire-based outcomes better than traditional structures derived from questionnaire data [65].

These novel latent structures are uncovered by leveraging the unique advantages of FMs in ways traditional surveys cannot. FMs can analyze data in dozens of languages, process vast numbers of items and lexical descriptions, examine historical texts, operate at extremely large scales with minimal cost, and access unprompted expressions of human psychology. Empirical findings also suggest that the structures identified through questionnaire data are neither the only valid representations nor necessarily the optimal. Furthermore, the ability of FMs to model constructs across multiple modalities opens new theoretical avenues for examining the coherence (or lack thereof) between an individual's self-perception, written expression, and behavior. Taken together, FMs can serve as a computational "microscope," revealing more nuanced semantics than previous computational models, or as a "telescope," enabling the analysis of data at scale and across cultures, languages, modalities, and time.

4.3 Modeling cognition and representing populations

FMs present an unprecedented opportunity to simulate individuals and populations (§ 2.2). This potential stems from two key strengths of FMs: their versatility in generative modeling, and their remarkable ability to learn from context and absorb complex patterns within vast training data.

Traditionally, the development of psychological and cognitive theories has followed a hypothesis-driven approach: researchers formulate theoretical predictions and test them in small, domain-specific experiments. The resulting data are then used to fit simple statistical models, which are refined iteratively. The process is often limited by the scale of available data and the narrow scope of experimental designs.

Recent FMs are opening new pathways for theory development [12, 109, 131, 169]. Cognitive modeling FMs (e.g., Centaur [12]) are trained on large-scale human experimental data and capture behavior across diverse domains. Such black-box models can serve as scaffolds for extracting interpretable theories. Recent work shows that reasoning-enhanced FMs (e.g., DeepSeek-R1 [44]) can be prompted to distill interpretable theories from human experimental data [12]. While these distilled theories may not match the predictive accuracy of Centaur, they already explain human behavior better than expert-proposed theories. Using black-box Centaur as references enables theory refinement without additional domain-specific data collection. This iterative process can produce interpretable theories that fit as well as Centaur [12]. Such FM-guided theory discovery challenges the common belief that black-box models must necessarily trade interpretability for accuracy [33].

Beyond modeling individual cognition, using FMs to represent populations has significant theoretical implications. Empirical findings from response simulation studies suggest that the FM training data are inherently biased samples of human discourse. Consequently, without specialized tuning, FMs most likely do not directly mirror individual minds but instead reflect collective narratives, stereotypes, and folk theories embedded in cultures and societies [18, 126]. This perspective underscores their unique theoretical value. Vanilla FMs can be powerful tools that enable researchers to model how societal biases are encoded in language [22], trace the evolution of values and moral frameworks across historical periods (by training on texts from different eras) [112], and examine how collective beliefs and attitudes are constructed and propagated [102]. For instance, large-scale analyses of nine language models and 1.4 million online images reveal systematic age-related gender stereotypes, with women consistently represented as younger than men across occupations—a distortion absent in actual workforce demographics [42]. Such findings demonstrate the theoretical value of FMs for exposing and quantifying how societal biases are encoded in digital discourse at an unprecedented scale.

4.4 LLM psychometrics

LLM psychometrics is referred to as an interdisciplinary field dedicated to evaluating, understanding, and enhancing LLMs through the application and integration of psychometric instruments, theories, and principles [157]. Studying LLMs as test subjects prompts a revisit of basic psychometric assumptions concerning subjects, constructs, and measurement methods.

Redefining the "subject" of measurement. Psychometrics was originally developed to assess humans. However, its recent and rapidly growing applications to FMs, especially LLMs, fundamentally challenge this core premise. Whereas human psychology assumes that traits such as personality are relatively stable over time, LLMs display contextually adaptive traits; they do not possess a single, enduring personality, but instead can adopt a wide range of personas depending on the prompts they receive [122]. There is a lack of consensus in how LLMs are treated in this line of research: a single LLM can be treated as an individual subject in its default state [158], as a population of subjects through diverse role-playing prompts [122], or as a monolithic entity generating stochastic, diverse responses to the same item [46]. This flexibility disrupts the traditional psychometric notion of a single test-taker with a unique set of traits. Additionally, the very act of applying psychometric tests to LLMs raises the issue of anthropomorphization. Attributing human-like consciousness or intentions to these models can mislead interpretation. As a result, psychometricians are compelled to a new theoretical frontier. Interdisciplinary efforts are needed to carefully distinguish between measuring human psychology and quantifying behavioral patterns in a non-human entity, and to design new theoretical systems for the latter as needed [157].

Re-examining the universality of psychological constructs. Research indicates that psychological constructs developed for humans may not be directly applicable to LLMs, compelling the field to consider how intelligence can be modeled beyond the human domain. Studies consistently show that human-centric theoretical models, such as the Big Five and Schwartz’s Value Theory, do not fit LLMs well [105, 136]. For instance, the clear separation of Big Five personality traits observed in humans is not reflected in LLM responses, and human value structures do not map cleanly onto LLM outputs. This mismatch necessitates a shift from relying solely on human constructs to developing constructs that are native to LLMs. Recent pioneering research addresses this gap by deriving new value and ability systems directly from the behavioral patterns of LLMs themselves [17, 159]. The newly established frameworks often indicate better validity, for example, in predicting and ensuring LLM safety [159]. Psychometrics may need to evolve from a science focused exclusively on human psychology to a broader discipline capable of measuring latent traits in any intelligent system. In doing so, theoretical models of these constructs may need to be reestablished.

Toward more ecologically valid measurement approaches. The unique properties of LLMs compel psychometrics to move beyond traditional, human-centric methods, particularly self-report inventories, toward more ecologically valid, behavior-based approaches. Unlike humans, LLMs lack subjective experience and have limited self-knowledge, so their responses to self-report items may be shaped more by statistical regularities than by introspection [130]. This leads to systematic artifacts such as option order effects [29], a preference for socially desirable answers [130], and sycophantic agreement regardless of the items [91, 136], all of which undermine the interpretability and validity of self-report data. Compounding these issues, LLMs are highly sensitive to subtle variations in prompt wording, syntax, and response format, resulting in inconsistent outputs that challenge the psychometric principle of parallel forms reliability [118]. These unique properties (or limitations in current AI training) necessitate alternative strategies that better capture the behavioral repertoire of LLMs. Approaches such as analyzing open-ended responses [117], observing agentic behavior in simulated environments [168], and leveraging naturalistic interactions [158] offer more ecologically and empirically valid results. In addition, applying these approaches to LLMs is much simpler and more cost-efficient than to humans, as AI is readily accessible anytime, anywhere, to any query. Taken together, we argue that while self-reports are the gold standard in human psychometrics, ecologically valid approaches should be the go-to choice in LLM psychometrics.

5 Recommendations

The integration of FMs into psychometrics is rapidly advancing, yet consensus on best practices remains limited. To guide this integration, we offer the following recommendations for researchers and practitioners.

5.1 Conceptual and methodological rigor

FMs are tools for prototyping and hypothesis generation, not yet ready replacements for human participants. FMs do not share the cognitive mechanisms of human psychology; they are complex statistical models trained on vast datasets. Currently, treating them as direct substitutes for human participants can lead to flawed conclusions [81]. Instead, their strength lies in augmenting the research process. Use FMs to rapidly prototype instruments, generate novel hypotheses, and explore theoretical boundaries in-silico before costly human data collection [41, 140]. A staged roadmap for LLM social simulations suggests using them for pilot and exploratory studies in the near term; their use could expand to exact replications and large-scale or counterfactual studies as methods improve [3].

Methodological advances should be accompanied by rigorous psychometric validation against human data. This process requires comprehensive psychometric validation that goes beyond simply comparing mean scores. For example, researchers may assess the factor structure, reliability, validity, and measurement invariance of AI-generated data. Early, head-to-head comparisons with human data are essential before scaling up generation and experiments. Iterative validation frameworks can help continuously track methodological progress and identify limitations systematically.

Consider enhancing FM performance for psychometric applications when out-of-the-box FMs are not sufficient. As detailed in § 3, use prompt engineering to control generation, scaffold complex

reasoning, and probe model behavior. When data is available, fine-tuning can adapt FMs to specific domains. For complex applications, consider agentic AI architectures that orchestrate multiple models or tools.

5.2 Robustness and reliability

Design for and test against perturbations and response sets. Findings based on FMs can be brittle. Research shows that FMs (1) can be sensitive to variations in prompts, languages, and task formulations [117, 118]; (2) exhibit response set biases, such as preferring option A in multiple-choice questions [29], favoring confirmatory responses [91], and giving generally high or low scores regardless of input [158]; and (3) are partially driven by social desirability [92, 130]. To ensure robust and reliable results, systematically compare outputs across multiple models and rigorously test sensitivity to artifacts. Use perturbation tests (e.g., paraphrasing prompts or altering response option order) to evaluate result stability. Assess self-consistency by generating and comparing multiple outputs from the same model. One should maintain vigilance for artifacts throughout study design and interpretation.

Address potential data contamination and the temporal context of FMs. FMs may have been exposed to evaluation tasks and data during training, which can artificially boost performance and undermine validity. Therefore, it is important to guard against data contamination to ensure robust findings. For example, when using FMs for experimental replication, it is ideal to use results from studies published after the model’s knowledge cutoff. Conversely, when leveraging LLMs’ encoded knowledge, one should recognize the temporal limitations imposed by their training data.

Ensure transparent and comprehensive reporting. Reproducibility is a cornerstone of scientific progress. When reporting results from FM-based research, provide complete documentation of your methods. This should include the models used, model versions, full prompts, hyperparameters such as decoding settings, and other implementation details. To support comprehensive reporting, [Box 1](#) offers a checklist that supplements traditional guidelines for psychological research by addressing considerations unique to FM-based methods.

5.3 Responsible and ethical integration

Make fairness a core principle. Ethical use of FMs in psychometrics requires a strong commitment to fairness. Researchers should conduct statistical analyses for differential item or algorithmic functioning [133]. Audits for bias should be implemented by evaluating FM-generated items and responses across diverse demographic groups to identify and mitigate biases [37, 137]. When necessary, apply bias mitigation strategies [137].

Foster interdisciplinary collaboration. Progress in this rapidly evolving field depends on close collaboration between AI experts and psychologists. AI researchers contribute technical knowledge, while psychologists ensure applications are grounded in sound measurement and ethics. Such partnerships are vital for developing tools that are reliable, valid, and fair.

Adopt a socio-technical perspective on governance. Treat FM-based psychometric tools as socio-technical systems when deploying them at scale (e.g., personality tests on websites). Carefully assess potential risks and implement holistic governance that accounts for societal impacts, following established standards [10].

Box 1 | Additional reporting checklist for FM-based psychometric research

Model information

- ☐ Model name(s), version(s), and release date(s)
- ☐ Model provider/source
- ☐ Model access method (API, local inference)
- ☐ Knowledge cutoff date
- ☐ Justification for model selection criteria

Experimental design

- ☐ Complete prompts and instructions (including system messages)
- ☐ Hyperparameters (temperature, top-p, max tokens, etc.)
- ☐ Number of runs/samples generated per condition

Validation and quality control

- ☐ Psychometric validation methods and results
- ☐ Mitigation strategies for data contamination
- ☐ Perturbation tests and sensitivity analyses
- ☐ Response quality filtering and exclusion criteria

Ethical and fairness considerations

- ☐ Demographic groups and characteristics examined
- ☐ Bias detection methods and statistical tests employed
- ☐ Bias mitigation strategies implemented (if applicable)
- ☐ Potential risks and harms associated with FM use
- ☐ Intended use and potential misuse considerations

Reproducibility

- ☐ Code availability
- ☐ Data availability (synthetic data, prompts, model outputs, etc.)
- ☐ Computational environment and dependencies
- ☐ Reproducibility measures
- ☐ Estimated computational costs and resources required

Limitations and scope

- ☐ Key limitations of the FM-based approach
- ☐ Scope and applicability of the findings

6 Summary and future directions

The integration of FMs into psychometrics promises a paradigm shift. It broadens the scope of what can be measured, introduces new methodologies for measurement, enables assessments across diverse settings, and reduces associated costs. Psychometrics-specific enhancement strategies further improve both the theoretical and empirical validity of this integration. On the theoretical front, FMs facilitate data-driven taxonomic clarification, the discovery of novel latent structures, streamlined theory development, and the emergence of a dedicated field of LLM psychometrics.

Research at the intersection of FMs and psychometrics is expanding rapidly, yet the field remains in its infancy. Several key scientific inquiries remain underdeveloped. First, future research should develop broader, enhanced, and more integrated applications of FMs in psychometrics. This includes designing agentic AI for adaptive testing and instrument creation, building multimodal interactive assessment systems for more engaging and valid assessments, and calibrating LLMs for individualized in-silico experiments. Second, the adaptation and enhancement of FMs for psychometric applications demands ongoing interdisciplinary collaboration. Current approaches are still limited in their generality and dependency on heavily annotated data. State-of-the-art AI methods (e.g., self-evolving algorithms) are yet to be leveraged for this purpose. Third, the theoretical impact of FMs in psychology requires further study. FMs could enable unified, data-driven taxonomies of psychological constructs and accelerate theory discovery by extracting interpretable models from large datasets. Fourth, LLM psychometrics is an emerging field that seeks rigorous measurement of analogous psychological traits

in LLMs. Interdisciplinary collaboration is needed to establish theoretical foundations, define AI-native constructs, develop valid behavior-based assessment methods, and explore algorithms to steer AI character. Finally, ethical and robust deployment of these tools is essential. Future work should develop scalable methods to ensure FM consistency and alignment, establish reliable and robust findings, and audit and mitigate algorithmic bias. Building a trustworthy science of FM-augmented psychometrics requires establishing standards, benchmarks, and platforms for transparency and reproducibility.

Collectively, the integration of FMs promises a transformative chapter for psychometrics. Realizing this potential will require interdisciplinary efforts to advance the methodological, theoretical, and ethical frontiers of the field.

A Terminologies

Box 2 | Key AI terminology

Foundation Models (FMs) — Large-scale, deep-learning models trained on vast amounts of data, which can be used for and adapted to a wide range of downstream tasks.

Large Language Models (LLMs) — A type of FM specifically trained on massive text datasets to understand, generate, and process human language. Examples include GPT-4 [1] and DeepSeek [44].

Vision Foundation Models (VFMs) — A type of FM specifically trained on massive datasets of images and videos to understand, generate, and process visual information.

Generative AI — A class of AI models, including many FMs, that can create new content, such as text, images, or code, rather than simply analyzing or classifying existing content.

Semantic embeddings — Numerical vector representations of words, sentences, or other content in a high-dimensional space. The geometric distance between vectors in this space reflects their semantic similarity, allowing models to capture nuanced meaning.

Zero-shot, few-shot, and in-context learning — The ability of FMs to perform tasks without any task-specific training (zero-shot) or with only a few examples provided in the prompt (few-shot). This learning happens "in-context" during a single interaction.

Prompt engineering — The practice of carefully designing input prompts to guide FMs toward producing desired outputs. It is a key technique for controlling and optimizing the behavior of these models without retraining them.

Retrieval-augmented generation (RAG) — A technique that enhances FMs by allowing them to retrieve information from an external knowledge base (e.g., a collection of documents) before generating a response. This helps to ground outputs in factual, up-to-date information.

Supervised learning — A fundamental machine learning paradigm where a model is trained on a dataset containing input-output pairs (labeled data). The model learns to map inputs to outputs, enabling it to make predictions on new, unseen data.

Reinforcement learning — A machine learning paradigm in which an agent interacts with an environment by taking actions and receiving feedback in the form of rewards or penalties. Through repeated trial and error, the agent learns a policy that maximizes cumulative reward over time.

Alignment — The process of ensuring that an AI system's goals and behaviors are consistent with human values and intentions. Techniques like Reinforcement Learning from Human Feedback (RLHF) are used to align models, making them more helpful, harmless, and honest. In particular, safety alignment ensures the model avoids harmful or biased outputs and adheres to ethical guidelines.

Fine-tuning — The process of further training a pretrained FM on a smaller, task-specific dataset. This adapts the model to a particular application, often improving its performance and alignment with desired behaviors.

Chain-of-Thought (CoT) prompting [148] — A prompt engineering technique that encourages an LLM to explain its reasoning process step-by-step before giving a final answer. This often improves performance on complex reasoning tasks by mimicking a more deliberate thought process.

LLM-as-a-Judge — A method where an LLM is prompted to act as an evaluator or "judge" to score, rate, or classify content based on a given set of criteria or rubrics.

Agentic AI — AI systems designed to be autonomous and goal-directed. They can plan, reason, use tools, and collaborate with other agents to accomplish complex tasks, often involving interaction with an external environment.

Multi-agent systems — A subfield of AI that involves the design and coordination of multiple autonomous agents to solve problems that are beyond the capabilities of a single agent.

Box 3 | Key psychometric terminology

Latent psychological constructs — Unobservable psychological attributes, traits, or abilities that are inferred from observable behaviors, such as responses to questionnaire items.

Reliability and validity — Two fundamental criteria for evaluating the quality of a measurement instrument. Reliability refers to the consistency and stability of measurement results over time, across items, or between raters. Validity refers to the extent to which the instrument accurately measures the construct it is intended to measure.

Instrument — A structured tool used to obtain observable indicators of a construct. In psychology, instruments include questionnaires, tests, or interview protocols comprising items designed to elicit responses that reflect underlying traits or states.

Item — A single, discrete prompt or question within an instrument that elicits an observable response indicative of the target construct. Items may be statements, questions, tasks, or stimuli with specified response formats (e.g., multiple choice, Likert agreement, open-ended).

Scale — A set of items with a specified response format (options and anchors) that maps responses to a measurement level (nominal, ordinal, interval, or ratio) for a target construct.

Item Response Theory (IRT) — A family of statistical models that link individual item responses to the underlying latent trait being measured. IRT models can estimate item properties such as difficulty (how much of the latent trait is needed to endorse the item) and discrimination (how well the item differentiates between individuals with different levels of the trait).

Factor structure — The underlying dimensional structure of a set of variables or test items. Factor analysis is used to identify common factors that explain the correlations among items, providing evidence for the internal structure of a construct.

Jingle-jangle fallacies — Two common errors in psychological research. The jingle fallacy is the mistaken belief that two different things are the same because they bear the same name. The jangle fallacy is the mistaken belief that two identical things are different because they are labeled differently.

Adaptive testing — A method of assessment in which the selection of items is dynamically adapted to the test-taker's performance level. Items are chosen to be maximally informative based on the individual's previous responses, making the test more efficient and precise.

Differential Item Functioning (DIF) — Occurs when an item on a test has different statistical properties for different subgroups of test-takers, after controlling for the underlying ability being measured. DIF is a source of test bias that can lead to unfairness.

Measurement invariance — A statistical property indicating that a measurement instrument functions in the same way across different groups (e.g., demographic or cultural groups). Establishing measurement invariance is crucial for making meaningful and fair comparisons of scores across these groups.

Internal consistency — A measure of reliability that assesses the extent to which items on a test that propose to measure the same construct produce similar scores. Cronbach's alpha is a common estimator.

Construct validity — The degree to which a test measures the theoretical construct it purports to measure.

Response sets — A tendency for individuals to respond to test items in a systematic way that is unrelated to the item content. Examples include acquiescence bias (yea-saying).

Social desirability bias — A systematic tendency to provide responses that present oneself favorably according to perceived social norms, regardless of true feelings or behaviors.

References

- [1] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023. URL <https://arxiv.org/abs/2303.08774>.
- [2] Gati V Aher, Rosa I. Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 337–371. PMLR, 2023. URL <https://proceedings.mlr.press/v202/aher23a.html>.
- [3] Jacy Reese Anthis, Ryan Liu, Sean M Richardson, Austin C Kozlowski, Bernard Koch, Erik Brynjolfsson, James Evans, and Michael S Bernstein. Position: Llm social simulations are a promising research method. In *Forty-second International Conference on Machine Learning Position Paper Track*, 2025. URL <https://openreview.net/forum?id=cRBg1dtj7o>.
- [4] Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, February 2023. ISSN 1476-4989. doi: 10.1017/pan.2023.2. URL <http://dx.doi.org/10.1017/pan.2023.2>.
- [5] Jan Ketil Arnulf, Kai Rune Larsen, Øyvind Lund Martinsen, and Chih How Bong. Predicting survey responses: How and why semantics shape survey statistics on organizational behaviour. *PLoS ONE*, 9(9):e106361, September 2014. ISSN 1932-6203. doi: 10.1371/journal.pone.0106361. URL <http://dx.doi.org/10.1371/journal.pone.0106361>.
- [6] Yigal Attali, Andrew Runge, Geoffrey T. LaFlair, Kevin Yancey, Sarah Goodwin, Yena Park, and Alina A. von Davier. The interactive reading task: Transformer-based automatic item generation. *Frontiers in Artificial Intelligence*, 5, July 2022. ISSN 2624-8212. doi: 10.3389/frai.2022.903077. URL <http://dx.doi.org/10.3389/frai.2022.903077>.
- [7] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, et al. Qwen2. 5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025.
- [8] Ryan S. J. d. Baker, Albert T. Corbett, and Vincent Aleven. *More Accurate Student Modeling through Contextual Estimation of Slip and Guess Probabilities in Bayesian Knowledge Tracing*, page 406–415. Springer Berlin Heidelberg, 2008. ISBN 9783540691327. doi: 10.1007/978-3-540-69132-7_44. URL http://dx.doi.org/10.1007/978-3-540-69132-7_44.
- [9] Joachim Baumann, Paul Röttger, Aleksandra Urman, Albert Wendsjö, Flor Miriam Plaza del Arco, Johannes B. Gruber, and Dirk Hovy. Large language model hacking: Quantifying the hidden risks of using llms for text annotation, 2025. URL <https://arxiv.org/abs/2509.08825>.
- [10] Yoshua Bengio, Geoffrey Hinton, Andrew Yao, Dawn Song, Pieter Abbeel, Trevor Darrell, Yuval Noah Harari, Ya-Qin Zhang, Lan Xue, Shai Shalev-Shwartz, Gillian Hadfield, Jeff Clune, Tegan Maharaj, Frank Hutter, Atılım Güneş Baydin, Sheila McIlraith, Qiqi Gao, Ashwin Acharya, David Krueger, Anca Dragan, Philip Torr, Stuart Russell, Daniel Kahneman, Jan Brauner, and Sören Mindermann. Managing extreme ai risks amid rapid progress. *Science*, 384(6698):842–845, May 2024. ISSN 1095-9203. doi: 10.1126/science.adn0117. URL <http://dx.doi.org/10.1126/science.adn0117>.
- [11] Giorgio Biancini, Alessio Ferrato, and Carla Limongelli. Multiple-choice question generation using large language models: Methodology and educator insights. In *Adjunct Proceedings of the 32nd ACM Conference on User Modeling, Adaptation and Personalization*, UMAP Adjunct '24, page 584–590, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400704666. doi: 10.1145/3631700.3665233. URL <https://doi.org/10.1145/3631700.3665233>.
- [12] Marcel Binz, Elif Akata, Matthias Bethge, Franziska Brändle, Fred Callaway, Julian Coda-Forno, Peter Dayan, Can Demircan, Maria K. Eckstein, Noémi Éltető, Thomas L. Griffiths, Susanne Haridi, Akshay K. Jagadish, Li Ji-An, Alexander Kipnis, Sreejan Kumar, Tobias Ludwig, Marvin Mathony, Marcelo Mattar, Alireza Modirshanechi, Surabhi S. Nath, Joshua C.

- Peterson, Milena Rmus, Evan M. Russek, Tankred Saanum, Johannes A. Schubert, Luca M. Schulze Buschoff, Nishad Singhi, Xin Sui, Mirko Thalmann, Fabian J. Theis, Vuong Truong, Vishaal Udandaraao, Konstantinos Voudouris, Robert Wilson, Kristin Witte, Shuchen Wu, Dirk U. Wulff, Huadong Xiong, and Eric Schulz. A foundation model to predict and capture human cognition. *Nature*, July 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-09215-4. URL <http://dx.doi.org/10.1038/s41586-025-09215-4>.
- [13] James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. Synthetic replacements for human survey data? the perils of large language models. *Political Analysis*, 32(4):401–416, May 2024. ISSN 1476-4989. doi: 10.1017/pan.2024.5. URL <http://dx.doi.org/10.1017/pan.2024.5>.
- [14] Rishi Bommasani, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, Erik Brynjolfsson, Shyamal Buch, Dallas Card, Rodrigo Castellon, Niladri Chatterji, Annie Chen, Kathleen Creel, Jared Quincy Davis, Dora Demszky, Chris Donahue, Moussa Doumbouya, Esin Durmus, Stefano Ermon, John Etchemendy, Kawin Ethayarajh, Li Fei-Fei, Chelsea Finn, Trevor Gale, Lauren Gillespie, Karan Goel, Noah Goodman, Shelby Grossman, Neel Guha, Tatsunori Hashimoto, Peter Henderson, John Hewitt, Daniel E. Ho, Jenny Hong, Kyle Hsu, Jing Huang, Thomas Icard, Saahil Jain, Dan Jurafsky, Pratyusha Kalluri, Siddharth Karamcheti, Geoff Keeling, Fereshte Khani, Omar Khattab, Pang Wei Koh, Mark Krass, Ranjay Krishna, Rohith Kudipudi, Ananya Kumar, Faisal Ladhak, Mina Lee, Tony Lee, Jure Leskovec, Isabelle Levent, Xiang Lisa Li, Xuechen Li, Tengyu Ma, Ali Malik, Christopher D. Manning, Suvir Mirchandani, Eric Mitchell, Zanele Munyikwa, Suraj Nair, Avanika Narayan, Deepak Narayanan, Ben Newman, Allen Nie, Juan Carlos Niebles, Hamed Nilforoshan, Julian Nyarko, Giray Ogut, Laurel Orr, Isabel Papadimitriou, Joon Sung Park, Chris Piech, Eva Portelance, Christopher Potts, Aditi Raghunathan, Rob Reich, Hongyu Ren, Frieda Rong, Yusuf Roohani, Camilo Ruiz, Jack Ryan, Christopher Ré, Dorsa Sadigh, Shiori Sagawa, Keshav Santhanam, Andy Shih, Krishnan Srinivasan, Alex Tamkin, Rohan Taori, Armin W. Thomas, Florian Tramèr, Rose E. Wang, William Wang, Bohan Wu, Jiajun Wu, Yuhuai Wu, Sang Michael Xie, Michihiro Yasunaga, Jiaxuan You, Matei Zaharia, Michael Zhang, Tianyi Zhang, Xikun Zhang, Yuhui Zhang, Lucia Zheng, Kaitlyn Zhou, and Percy Liang. On the opportunities and risks of foundation models, 2022. URL <https://arxiv.org/abs/2108.07258>.
- [15] Denny Borsboom. Measuring the mind: Conceptual issues in contemporary psychometrics. 2005. URL <https://psycnet.apa.org/record/2005-16064-000>.
- [16] Ryan L Boyd, Ashwini Ashokkumar, Sarah Seraj, and James W Pennebaker. The development and psychometric properties of liwc-22. *Austin, TX: University of Texas at Austin*, 10:1–47, 2022. URL https://www.liwc.net/wp-content/uploads/2022/03/LIWC22_Manual.pdf.
- [17] Ryan Burnell, Han Hao, Andrew R. A. Conway, and Jose Hernandez Orallo. Revealing the structure of language model capabilities, 2023. URL <https://arxiv.org/abs/2306.10062>.
- [18] Jason W. Burton, Ezequiel Lopez-Lopez, Shahar Hechtlinger, Zoe Rahwan, Samuel Aeschbach, Michiel A. Bakker, Joshua A. Becker, Aleks Berdichevskaia, Julian Berger, Levin Brinkmann, Lucie Flek, Stefan M. Herzog, Saffron Huang, Sayash Kapoor, Arvind Narayanan, Anne-Marie Nussberger, Taha Yasseri, Pietro Nickl, Abdullah Almaatouq, Ulrike Hahn, Ralf H. J. M. Kurvers, Susan Leavy, Iyad Rahwan, Divya Siddarth, Alice Siu, Anita W. Woolley, Dirk U. Wulff, and Ralph Hertwig. How large language models can reshape collective intelligence. *Nature Human Behaviour*, 8(9):1643–1655, September 2024. ISSN 2397-3374. doi: 10.1038/s41562-024-01959-9. URL <http://dx.doi.org/10.1038/s41562-024-01959-9>.
- [19] Matthew Byrd and Shashank Srivastava. Predicting difficulty and discrimination of natural language questions. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio, editors, *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 119–130, Dublin, Ireland, May 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-short.15. URL <https://aclanthology.org/2022.acl-short.15/>.
- [20] Yong Cao, Haijiang Liu, Arnav Arora, Isabelle Augenstein, Paul Röttger, and Daniel Herscovich. Specializing large language models to simulate survey response distributions

- for global populations. In Luis Chiruzzo, Alan Ritter, and Lu Wang, editors, *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3141–3154, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics. ISBN 979-8-89176-189-6. doi: 10.18653/v1/2025.naacl-long.162. URL <https://aclanthology.org/2025.naacl-long.162/>.
- [21] Cheng Cheng, GuanHao Zhao, Zhenya Huang, Yan Zhuang, Zhaoyuan Pan, Qi Liu, Xin Li, and Enhong Chen. Towards explainable computerized adaptive testing with large language model. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 2655–2672, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-emnlp.149. URL <https://aclanthology.org/2024.findings-emnlp.149/>.
 - [22] Vanessa Cheung, Maximilian Maier, and Falk Lieder. Large language models show amplified cognitive biases in moral decision-making. *Proceedings of the National Academy of Sciences*, 122(25), June 2025. ISSN 1091-6490. doi: 10.1073/pnas.2412015122. URL <http://dx.doi.org/10.1073/pnas.2412015122>.
 - [23] Yu Ying Chiu, Liwei Jiang, and Yejin Choi. Dailydilemmas: Revealing value preferences of llms with quandaries of daily life. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=PGhiPGBf47>.
 - [24] Yun-Shiuan Chuang, Krirk Nirunwiroj, Zach Studdiford, Agam Goyal, Vincent Frigo, Sijia Yang, Dhavan Shah, Junjie Hu, and Timothy Rogers. Beyond demographics: Aligning role-playing llm-based agents using human belief networks. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 14010–14026, 2024.
 - [25] Ziyang Cui, Ning Li, and Huaikang Zhou. A large-scale replication of scenario-based experiments in psychology and management using large language models. *Nature Computational Science*, 5(8):627–634, July 2025. ISSN 2662-8457. doi: 10.1038/s43588-025-00840-7. URL <http://dx.doi.org/10.1038/s43588-025-00840-7>.
 - [26] Andrew Cutler and David M. Condon. Deep lexical hypothesis: Identifying personality structure in natural language. *Journal of Personality and Social Psychology*, 125(1):173–197, July 2023. ISSN 0022-3514. doi: 10.1037/pspp0000443. URL <http://dx.doi.org/10.1037/pspp0000443>.
 - [27] Arthur P Dempster, Nan M Laird, and Donald B Rubin. Maximum likelihood from incomplete data via the em algorithm. *Journal of the royal statistical society: series B (methodological)*, 39(1):1–22, 1977. URL <https://rss.onlinelibrary.wiley.com/doi/abs/10.1111/j.2517-6161.1977.tb01600.x>.
 - [28] Danica Dillion, Niket Tandon, Yuling Gu, and Kurt Gray. Can ai language models replace human participants? *Trends in Cognitive Sciences*, 27(7):597–600, July 2023. ISSN 1364-6613. doi: 10.1016/j.tics.2023.04.008. URL <http://dx.doi.org/10.1016/j.tics.2023.04.008>.
 - [29] Ricardo Dominguez-Olmedo, Moritz Hardt, and Celestine Mendler-Dünnér. Questioning the survey responses of large language models. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 45850–45878. Curran Associates, Inc., 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/515c62809e0a29729d7eec26e2916fc0-Paper-Conference.pdf.
 - [30] Bangde Du, Ziyi Ye, Zhijing Wu, Jankowska Monika, Shuqi Zhu, Qingyao Ai, Yujia Zhou, and Yiqun Liu. Valuesim: Generating backstories to model individual value systems, 2025. URL <https://arxiv.org/abs/2505.23827>.
 - [31] Fang, Qixiang, Nguyen, Dong, and Oberski, Daniel L. Evaluating the construct validity of text embeddings with application to survey questions. *EPJ Data Sci.*, 11(1):39, 2022. doi: 10.1140/epjds/s13688-022-00353-7. URL <https://doi.org/10.1140/epjds/s13688-022-00353-7>.
 - [32] Tommaso Feraco and Enrico Toffalini. Sembeddings: how to evaluate model misfit before data collection using large-language models. *Frontiers in Psychology*, 15, February 2025. ISSN 1664-1078. doi: 10.3389/fpsyg.2024.1433339. URL <http://dx.doi.org/10.3389/fpsyg.2024.1433339>.

- [33] Stefan Feuerriegel, Abdurahman Maarouf, Dominik Bär, Dominique Geissler, Jonas Schweisthal, Nicolas Pröllochs, Claire E. Robertson, Steve Rathje, Jochen Hartmann, Saif M. Mohammad, Oded Netzer, Alexandra A. Siegel, Barbara Plank, and Jay J. Van Bavel. Using natural language processing to analyse text data in behavioural science. *Nature Reviews Psychology*, 4(2):96–111, January 2025. ISSN 2731-0574. doi: 10.1038/s44159-024-00392-z. URL <http://dx.doi.org/10.1038/s44159-024-00392-z>.
- [34] Ricardo Flores, ML Tlachac, Avantika Shrestha, and Elke A. Rundensteiner. Wavface: A multimodal transformer-based model for depression screening. *IEEE Journal of Biomedical and Health Informatics*, 29(5):3632–3641, 2025. doi: 10.1109/JBHI.2025.3529348.
- [35] Yuan Gao, Dokyun Lee, Gordon Burtch, and Sina Fazelpour. Take caution in using llms as human surrogates. *Proceedings of the National Academy of Sciences*, 122(24), June 2025. ISSN 1091-6490. doi: 10.1073/pnas.2501660122. URL <http://dx.doi.org/10.1073/pnas.2501660122>.
- [36] Mark J. Gierl and Thomas M. Haladyna. *Automatic Item Generation*. Routledge, August 2012. ISBN 9781136636899. doi: 10.4324/9780203803912. URL <http://dx.doi.org/10.4324/9780203803912>.
- [37] Moshe Glickman and Tali Sharot. How human–ai feedback loops alter human perceptual, emotional and social judgements. *Nature Human Behaviour*, 9(2):345–359, December 2024. ISSN 2397-3374. doi: 10.1038/s41562-024-02077-2. URL <http://dx.doi.org/10.1038/s41562-024-02077-2>.
- [38] Lewis R Goldberg. An alternative “description of personality”: The big-five factor structure. In *Personality and personality disorders*, pages 34–47. Routledge, 2013. URL <https://www.scis.com/chapters/edit/10.4324/9780203822845-4/alternative-description-personality-big-five-factor-structure-lewis-goldberg>.
- [39] Lewis R. Goldberg, John A. Johnson, Herbert W. Eber, Robert Hogan, Michael C. Ashton, C. Robert Cloninger, and Harrison G. Gough. The international personality item pool and the future of public-domain personality measures. *Journal of Research in Personality*, 40(1):84–96, February 2006. ISSN 0092-6566. doi: 10.1016/j.jrp.2005.08.007. URL <http://dx.doi.org/10.1016/j.jrp.2005.08.007>.
- [40] Guher Gorgun and Okan Bulut. Instruction-tuned large-language models for quality control in automatic item generation: A feasibility study. *Educational Measurement: Issues and Practice*, 44(1):96–107, December 2024. ISSN 1745-3992. doi: 10.1111/emip.12663. URL <http://dx.doi.org/10.1111/emip.12663>.
- [41] Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, Tao Tu, Anil Palepu, Petar Sirkovic, Artiom Myaskovsky, Felix Weissenberger, Keran Rong, Ryutaro Tanno, Khaled Saab, Dan Popovici, Jacob Blum, Fan Zhang, Katherine Chou, Avinatan Hassidim, Burak Gokturk, Amin Vahdat, Pushmeet Kohli, Yossi Matias, Andrew Carroll, Kavita Kulkarni, Nenad Tomasev, Yuan Guan, Vikram Dhillon, Eeshit Dhaval Vaishnav, Byron Lee, Tiago R D Costa, José R Penadés, Gary Peltz, Yunhan Xu, Annalisa Pawlosky, Alan Karthikesalingam, and Vivek Natarajan. Towards an ai co-scientist, 2025. URL <https://arxiv.org/abs/2502.18864>.
- [42] Douglas Guilbeault, Solène Delecourt, and Bhargav Srinivasa Desikan. Age and gender distortion in online media and large language models. *Nature*, October 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-09581-z. URL <http://dx.doi.org/10.1038/s41586-025-09581-z>.
- [43] Gauthier Guinet, Behrooz Omidvar-Tehrani, Anoop Deoras, and Laurent Callot. Automated evaluation of retrieval-augmented language models with task-specific exam generation. In *Proceedings of the 41st International Conference on Machine Learning, ICML’24*. JMLR.org, 2024.
- [44] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Hanwei Xu, Honghui Ding, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jingchang Chen, Jingyang Yuan, Jinhao Tu, Junjie Qiu, Junlong Li, J. L. Cai,

- Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaichao You, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Mingxu Zhou, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiushi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, S. S. Li, Shuang Zhou, Shaoqing Wu, Tao Yun, Tian Pei, Tianyu Sun, T. Wang, Wangding Zeng, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, W. L. Xiao, Wei An, Xiaodong Liu, Xiaohan Wang, Xiaokang Chen, Xiaotao Nie, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, X. Q. Li, Xiangyue Jin, Xiaojin Shen, Xiaosha Chen, Xiaowen Sun, Xiaoxiang Wang, Xinnan Song, Xinyi Zhou, Xianzu Wang, Xinxia Shan, Y. K. Li, Y. Q. Wang, Y. X. Wei, Yang Zhang, Yanhong Xu, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Wang, Yi Yu, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yuan Ou, Yuduan Wang, Yue Gong, Yuheng Zou, Yujia He, Yunfan Xiong, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Y. X. Zhu, Yanping Huang, Yaohui Li, Yi Zheng, Yuchen Zhu, Yunxian Ma, Ying Tang, Yukun Zha, Yuting Yan, Z. Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhicheng Ma, Zhigang Yan, Zhiyu Wu, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Zizheng Pan, Zhen Huang, Zhipeng Xu, Zhongyu Zhang, and Zhen Zhang. Deepseek-r1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645 (8081):633–638, September 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-09422-z. URL <http://dx.doi.org/10.1038/s41586-025-09422-z>.
- [45] Otsu Haavisto and Robin Welsch. Questionnaires for everyone: Streamlining cross-cultural questionnaire adaptation with gpt-based translation quality evaluation, 2024. URL <https://arxiv.org/abs/2407.20608>.
- [46] Thilo Hagendorff, Ishita Dasgupta, Marcel Binz, Stephanie C. Y. Chan, Andrew Lampinen, Jane X. Wang, Zeynep Akata, and Eric Schulz. Machine psychology, 2024. URL <https://arxiv.org/abs/2303.13988>.
- [47] Ivan Hernandez and Weiwen Nie. The ai-ip: Minimizing the guesswork of personality scale item development through artificial intelligence. *Personnel Psychology*, 76(4):1011–1035, October 2022. ISSN 1744-6570. doi: 10.1111/peps.12543. URL <http://dx.doi.org/10.1111/peps.12543>.
- [48] Luke Hewitt, Ashwini Ashokkumar, Isaias Ghezze, and Robb Willer. Predicting results of social science experiments using large language models. *preprint*, 2024. URL <https://samim.io/dl/Predicting%20results%20of%20social%20science%20experiments%20using%20large%20language%20models.pdf>.
- [49] Björn E. Hommel, Franz-Josef M. Wollang, Veronika Kotova, Hannes Zacher, and Stefan C. Schmukle. Transformer-based deep neural language modeling for construct-specific automatic item generation. *Psychometrika*, 87(2):749–772, June 2022. ISSN 1860-0980. doi: 10.1007/s11336-021-09823-9. URL <http://dx.doi.org/10.1007/s11336-021-09823-9>.
- [50] Jinpeng Hu, Ao Wang, Qianqian Xie, Hui Ma, Zhuo Li, and Dan Guo. Agentmental: An interactive multi-agent framework for explainable and adaptive mental health assessment, 2025. URL <https://arxiv.org/abs/2508.11567>.
- [51] Tiancheng Hu and Nigel Collier. Quantifying the persona effect in llm simulations, 2024. URL <https://arxiv.org/abs/2402.10811>.
- [52] Zhen Huang, Yitian Long, Kaiping Peng, and Song Tong. An embedding-based semantic analysis approach: A preliminary study on redundancy detection in psychological concepts operationalized by scales. *Journal of Intelligence*, 13(1):11, January 2025. ISSN 2079-3200. doi: 10.3390/jintelligence13010011. URL <http://dx.doi.org/10.3390/jintelligence13010011>.
- [53] EunJeong Hwang, Bodhisattwa Majumder, and Niket Tandon. Aligning language models to user opinions. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 5906–5919, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.393. URL <https://aclanthology.org/2023.findings-emnlp.393/>.

- [54] Calvin Isley, Joshua Gilbert, Evangelos Kassos, Michaela Kocher, Allen Nie, Emma Brunskill, Ben Domingue, Jake Hofman, Joscha Legewie, Teddy Svoronos, Charlotte Tuminelli, and Sharad Goel. Assessing the quality of ai-generated exams: A large-scale field study, 2025. URL <https://arxiv.org/abs/2508.08314>.
- [55] Yoshee Jain, John Hollandier, Amber He, Sunny Tang, Liang Zhang, and John Sabatini. *Exploring the Potential of Large Language Models for Estimating the Reading Comprehension Question Difficulty*, page 202–213. Springer Nature Switzerland, 2025. ISBN 9783031929670. doi: 10.1007/978-3-031-92967-0_14. URL http://dx.doi.org/10.1007/978-3-031-92967-0_14.
- [56] Guangyuan Jiang, Manjie Xu, Song-Chun Zhu, Wenjuan Han, Chi Zhang, and Yixin Zhu. Evaluating and inducing personality in pre-trained language models. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 10622–10643. Curran Associates, Inc., 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/21f7b745f73ce0d1f9bcea7f40b1388e-Paper-Conference.pdf.
- [57] Han Jiang, Xiaoyuan Yi, Zhihua Wei, Ziang Xiao, Shu Wang, and Xing Xie. Raising the bar: Investigating the values of large language models via generative evolving testing. In *Forty-second International Conference on Machine Learning*, 2024. URL <https://openreview.net/forum?id=OREM9ydeLZ>.
- [58] Liwei Jiang, Taylor Sorensen, Sydney Levine, and Yejin Choi. Can language models reason about individualistic human values and preferences?, 2025. URL <https://arxiv.org/abs/2410.03868>.
- [59] Nani Jin, Renjia Ye, and Peng Li. Diagnosis of depression based on facial multimodal data. *Frontiers in Psychiatry*, 16, January 2025. ISSN 1664-0640. doi: 10.3389/fpsy.2025.1508772. URL <http://dx.doi.org/10.3389/fpsy.2025.1508772>.
- [60] Antonio Laverghetta Jr., Simone Luchini, Averie Linell, Roni Reiter-Palmon, and Roger Beaty. The creative psychometric item generator: a framework for item generation and validation using large language models, 2024. URL <https://arxiv.org/abs/2409.00202>.
- [61] Radhika Kapoor, Sang T. Truong, Nick Haber, Maria Araceli Ruiz-Primo, and Benjamin W. Domingue. Prediction of item difficulty for reading comprehension items by creation of annotated item repository, 2025. URL <https://arxiv.org/abs/2502.20663>.
- [62] Elma Kerz, Yu Qiao, Sourabh Zanwar, and Daniel Wiechmann. Pushing on personality detection from verbal behavior: A transformer meets text contours of psycholinguistic features. In Jeremy Barnes, Orphée De Clercq, Valentin Barriere, Shabnam Tafreshi, Sawsan Alqahtani, João Sedoc, Roman Klinger, and Alexandra Balahur, editors, *Proceedings of the 12th Workshop on Computational Approaches to Subjectivity, Sentiment & Social Media Analysis*, pages 182–194, Dublin, Ireland, May 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.wassa-1.17. URL <https://aclanthology.org/2022.wassa-1.17/>.
- [63] Sevilay Kilmen and Okan Bulut. Shortening psychological scales: Semantic similarity matters. *Educational and Psychological Measurement*, February 2025. ISSN 1552-3888. doi: 10.1177/00131644251319047. URL <http://dx.doi.org/10.1177/00131644251319047>.
- [64] Michal Kosinski, David Stillwell, and Thore Graepel. Private traits and attributes are predictable from digital records of human behavior. *Proceedings of the National Academy of Sciences*, 110(15):5802–5805, March 2013. ISSN 1091-6490. doi: 10.1073/pnas.1218772110. URL <http://dx.doi.org/10.1073/pnas.1218772110>.
- [65] Vivek Kulkarni, Margaret L. Kern, David Stillwell, Michal Kosinski, Sandra Matz, Lyle Ungar, Steven Skiena, and H. Andrew Schwartz. Latent human traits in the language of social media: An open-vocabulary approach. *PLOS ONE*, 13(11):e0201703, November 2018. ISSN 1932-6203. doi: 10.1371/journal.pone.0201703. URL <http://dx.doi.org/10.1371/journal.pone.0201703>.
- [66] Jonas R. Kunst and Kinga Bierwiazzonek. Utilizing ai questionnaire translations in cross-cultural and intercultural research: Insights and recommendations. *International Journal of Intercultural Relations*, 97:101888, 2023. ISSN 0147-1767. doi: <https://doi.org/10.1016/j.ijintrel.2023.101888>. URL <https://www.sciencedirect.com/science/article/pii/S0147176723001360>.

- [67] Lucio La Cava and Andrea Tagarelli. Open models, closed minds? on agents capabilities in mimicking human personalities through open large language models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(2):1355–1363, 2025. doi: 10.1609/aaai.v39i2.32125. URL <https://ojs.aaai.org/index.php/AAAI/article/view/32125>.
- [68] Tin Lai, Yukun Shi, Zicong Du, Jiajie Wu, Ken Fu, Yichao Dou, and Ziqi Wang. Psy-llm: Scaling up global mental health psychological services with ai-based large language models, 2023. URL <https://arxiv.org/abs/2307.11991>.
- [69] John P. Lalor, Pedro Rodriguez, João Sedoc, and Jose Hernandez-Orallo. Item response theory for natural language processing. In Mohsen Mesgar and Sharid Loáiciga, editors, *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics: Tutorial Abstracts*, pages 9–13, St. Julian’s, Malta, March 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.eacl-tutorials.2. URL <https://aclanthology.org/2024.eacl-tutorials.2/>.
- [70] Kai R. Larsen and Chih How Bong. A tool for addressing construct identity in literature reviews and meta-analyses. *MIS Quarterly*, 40(3):529–552, 2016. ISSN 02767783, 21629730. URL <https://www.jstor.org/stable/26629026>.
- [71] Antonio Laverghetta Jr. and John Licato. Generating better items for cognitive assessments using large language models. In Ekaterina Kochmar, Jill Burstein, Andrea Horbach, Ronja Laarmann-Quante, Nitin Madnani, Anaïs Tack, Victoria Yaneva, Zheng Yuan, and Torsten Zesch, editors, *Proceedings of the 18th Workshop on Innovative Use of NLP for Building Educational Applications (BEA 2023)*, pages 414–428, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.bea-1.34. URL <https://aclanthology.org/2023.bea-1.34/>.
- [72] Esko Lehtonen, Tommi Buder-Gröndahl, and Sina Nordhoff. Revealing the influence of semantic similarity on survey responses: A synthetic data generation approach. *IEEE Access*, 13:40285–40301, 2025. ISSN 2169-3536. doi: 10.1109/access.2025.3546565. URL <http://dx.doi.org/10.1109/ACCESS.2025.3546565>.
- [73] Chang-Jin Li, Jiyan Zhang, Yun Tang, and Jian Li. Automatic item generation for personality situational judgment tests with large language models, 2025. URL <https://arxiv.org/abs/2412.12144>.
- [74] Cheng Li, Mengzhuo Chen, Jindong Wang, Sunayana Sitaram, and Xing Xie. Culturellm: Incorporating cultural differences into large language models. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 84799–84838. Curran Associates, Inc., 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/9a16935bf54c4af233e25d998b7f4a2c-Paper-Conference.pdf.
- [75] Jia Li, Yichao He, Jiacheng Xu, Tianhao Luo, Zhenzhen Hu, Richang Hong, and Meng Wang. Traits run deep: Enhancing personality assessment via psychology-guided llm representations and multimodal apparent behaviors, 2025. URL <https://arxiv.org/abs/2507.22367>.
- [76] Jingxuan Li, Yuning Yang, Shengqi Yang, Linfan Zhang, and Ying Nian Wu. Value-spectrum: Quantifying preferences of vision-language models via value decomposition in social media contexts, 2025. URL <https://arxiv.org/abs/2411.11479>.
- [77] Ming Li, Hong Jiao, Tianyi Zhou, Nan Zhang, Sydney Peters, and Robert W. Lissitz. Item difficulty modeling using fine-tuned small and large language models. *Educational and Psychological Measurement*, July 2025. ISSN 1552-3888. doi: 10.1177/00131644251344973. URL <http://dx.doi.org/10.1177/00131644251344973>.
- [78] Xueyan Li, Xinyan Chen, Yazhe Niu, Shuai Hu, and Yu Liu. Psydi: Towards a personalized and progressively in-depth chatbot for psychological measurements, 2025. URL <https://arxiv.org/abs/2408.03337>.
- [79] Zheng Li, Sujian Li, Dawei Zhu, Qilong Ma, and Weimin Xiong. EERPD: Leveraging emotion and emotion regulation for improving personality detection. In Owen Rambow, Leo Wanner, Marianna Apidianaki, Hend Al-Khalifa, Barbara Di Eugenio, and Steven Schockaert, editors, *Proceedings of the 31st International Conference on Computational Linguistics*, pages 7721–7734, Abu Dhabi, UAE, January 2025. Association for Computational Linguistics. URL <https://aclanthology.org/2025.coling-main.516/>.

- [80] Zhicheng Lin. Large language models as psychological simulators: A methodological guide, 2025. URL <https://arxiv.org/abs/2506.16702>.
- [81] Zhicheng Lin. Six fallacies in substituting large language models for human participants. *Advances in Methods and Practices in Psychological Science*, 8(3): 25152459251357566, 2025. doi: 10.1177/25152459251357566. URL <https://doi.org/10.1177/25152459251357566>.
- [82] Shuyu Liu, Ruoxi Wang, Ling Zhang, Xuequan Zhu, Rui Yang, Xinzhu Zhou, Fei Wu, Zhi Yang, Cheng Jin, and Gang Wang. Psychbench: A comprehensive and professional benchmark for evaluating the performance of llm-assisted psychiatric clinical practice, 2025. URL <https://arxiv.org/abs/2503.01903>.
- [83] Yunting Liu, Shreya Bhandari, and Zachary A. Pardos. Leveraging llm-respondents for item evaluation: a psychometric analysis, 2024. URL <https://arxiv.org/abs/2407.10899>.
- [84] Yunting Liu, Shreya Bhandari, and Zachary A. Pardos. Leveraging llm respondents for item evaluation: A psychometric analysis. *British Journal of Educational Technology*, 56(3):1028–1052, February 2025. ISSN 1467-8535. doi: 10.1111/bjet.13570. URL <http://dx.doi.org/10.1111/bjet.13570>.
- [85] Frederic M Lord and Melvin R Novick. *Statistical theories of mental test scores*. IAP, 2008. URL <https://psycnet.apa.org/record/1968-35040-000>.
- [86] Yuqing Lu, Xin Hu, and Dan Zhang. A prompt engineering method for generating emotional images for psychological research. *Affective Science*, July 2025. ISSN 2662-205X. doi: 10.1007/s42761-025-00315-y. URL <http://dx.doi.org/10.1007/s42761-025-00315-y>.
- [87] Yuxing Lu and Jinzhuo Wang. Karma: Leveraging multi-agent llms for automated knowledge graph enrichment. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL <https://arxiv.org/abs/2502.06472>.
- [88] Yuxing Lu, Xiaohong Liu, Zongxin Du, Yuanxu Gao, and Guangyu Wang. Medkpl: A heterogeneous knowledge enhanced prompt learning framework for transferable diagnosis. *Journal of Biomedical Informatics*, 143:104417, 2023. ISSN 1532-0464. doi: <https://doi.org/10.1016/j.jbi.2023.104417>. URL <https://www.sciencedirect.com/science/article/pii/S1532046423001387>.
- [89] Hotaka Maeda. Field-testing multiple-choice questions with ai examinees: English grammar items. *Educational and Psychological Measurement*, 85(2):221–244, October 2024. ISSN 1552-3888. doi: 10.1177/00131644241281053. URL <http://dx.doi.org/10.1177/00131644241281053>.
- [90] Julina Maharjan, Ruoming Jin, Jianfeng Zhu, and Deric Kenne. Psychometric evaluation of large language model embeddings for personality trait prediction. *J Med Internet Res*, 27:e75347, 2025. ISSN 1438-8871. doi: 10.2196/75347. URL <https://www.jmir.org/2025/1/e75347>.
- [91] Lars Malmqvist. Sycophancy in large language models: Causes and mitigations. In Kohei Arai, editor, *Intelligent Computing*, pages 61–74, Cham, 2025. Springer Nature Switzerland. ISBN 978-3-031-92611-2.
- [92] Mantas Mazeika, Xuwang Yin, Rishub Tamirisa, Jaehyuk Lim, Bruce W. Lee, Richard Ren, Long Phan, Norman Mu, Adam Khoja, Oliver Zhang, and Dan Hendrycks. Utility engineering: Analyzing and controlling emergent value systems in ais, 2025. URL <https://arxiv.org/abs/2502.08640>.
- [93] Arya D. McCarthy, Kevin P. Yancey, Geoffrey T. LaFlair, Jesse Egbert, Manqian Liao, and Burr Settles. Jump-starting item parameters for adaptive language tests. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 883–899, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.67. URL <https://aclanthology.org/2021.emnlp-main.67/>.
- [94] Lingrui Mei, Jiayu Yao, Yuyao Ge, Yiwei Wang, Baolong Bi, Yujun Cai, Jiazhi Liu, Mingyu Li, Zhong-Zhi Li, Duzhen Zhang, Chenlin Zhou, Jiayi Mao, Tianze Xia, Jiafeng Guo, and Shenghua Liu. A survey of context engineering for large language models, 2025. URL <https://arxiv.org/abs/2507.13334>.

- [95] Roy Meissner, Alexander Pögel, Katja Ihsberner, Martin Grüttmüller, Silvana Tornack, Andreas Thor, Norbert Pengel, Heinz-Werner Wollersheim, and Wolfram Hardt. Llm-generated competence-based e-assessment items for higher education mathematics: methodology and evaluation. *Frontiers in Education*, 9, October 2024. ISSN 2504-284X. doi: 10.3389/feduc.2024.1427502. URL <http://dx.doi.org/10.3389/feduc.2024.1427502>.
- [96] Samuel Messick. Validity of psychological assessment: Validation of inferences from persons' responses and performances as scientific inquiry into score meaning. *American Psychologist*, 50(9):741–749, September 1995. ISSN 0003-066X. doi: 10.1037/0003-066x.50.9.741. URL <http://dx.doi.org/10.1037/0003-066x.50.9.741>.
- [97] Nicola Milano, Maria Luongo, Michela Ponticorvo, and Davide Marocco. Semantic analysis of test items through large language model embeddings predicts a-priori factorial structure of personality tests. *Current Research in Behavioral Sciences*, 8:100168, 2025. ISSN 2666-5182. doi: 10.1016/j.crbeha.2025.100168. URL <http://dx.doi.org/10.1016/j.crbeha.2025.100168>.
- [98] Nicola Milano, Michela Ponticorvo, and Davide Marocco. Comparing human expertise and large language models embeddings in content validity assessment of personality tests, 2025. URL <https://arxiv.org/abs/2503.12080>.
- [99] Vlatko Nikolovski, Dimitar Trajanov, and Ivan Chorbev. Advancing ai in higher education: A comparative study of large language model-based agents for exam question generation, improvement, and evaluation. *Algorithms*, 18(3):144, March 2025. ISSN 1999-4893. doi: 10.3390/a18030144. URL <http://dx.doi.org/10.3390/a18030144>.
- [100] Alexander Novikov, Ngân Vū, Marvin Eisenberger, Emilien Dupont, Po-Sen Huang, Adam Zsolt Wagner, Sergey Shirobokov, Borislav Kozlovskii, Francisco J. R. Ruiz, Abbas Mehrabian, M. Pawan Kumar, Abigail See, Swarat Chaudhuri, George Holland, Alex Davies, Sebastian Nowozin, Pushmeet Kohli, and Matej Balog. Alphaevolve: A coding agent for scientific and algorithmic discovery, 2025. URL <https://arxiv.org/abs/2506.13131>.
- [101] Open-Source Psychometrics Project. Open-source psychometrics project, 2025. URL <https://openpsychometrics.org>. Accessed: 2025.
- [102] Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, UIST '23, page 1–22. ACM, October 2023. doi: 10.1145/3586183.3606763. URL <http://dx.doi.org/10.1145/3586183.3606763>.
- [103] Joon Sung Park, Carolyn Q. Zou, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Robb Willer, Percy Liang, and Michael S. Bernstein. Generative agent simulations of 1,000 people, 2024. URL <https://arxiv.org/abs/2411.10109>.
- [104] Avinash Patil and Amardeep Kour Gedhu. Cognitive-mental-llm: Evaluating reasoning in large language models for mental health prediction via online text, 2025. URL <https://arxiv.org/abs/2503.10095>.
- [105] Sanne Peereboom, Inga Schwabe, and Bennett Kleinberg. Cognitive phantoms in large language models through the lens of latent variables. *Computers in Human Behavior: Artificial Humans*, 4:100161, May 2025. ISSN 2949-8821. doi: 10.1016/j.chbah.2025.100161. URL <http://dx.doi.org/10.1016/j.chbah.2025.100161>.
- [106] Heinrich Peters and Sandra C Matz. Large language models can infer psychological dispositions of social media users. *PNAS Nexus*, 3(6), May 2024. ISSN 2752-6542. doi: 10.1093/pnasnexus/pgae231. URL <http://dx.doi.org/10.1093/pnasnexus/pgae231>.
- [107] Heinrich Peters, Moran Cerf, and Sandra C. Matz. Large language models can infer personality from free-form user interactions, 2024. URL <https://arxiv.org/abs/2405.13052>.
- [108] Nikolay B Petrov, Gregory Serapio-García, and Jason Rentfrow. Limited ability of llms to simulate human psychological behaviours: a psychometric analysis, 2024. URL <https://arxiv.org/abs/2405.07248>.
- [109] Ori Plonsky, Reut Apel, Eyal Ert, Moshe Tennenholtz, David Bourgin, Joshua C. Peterson, Daniel Reichman, Thomas L. Griffiths, Stuart J. Russell, Even C. Carter, James F. Cavanagh, and Ido Erev. Predicting human decisions with behavioural theories and machine learning.

- Nature Human Behaviour*, July 2025. ISSN 2397-3374. doi: 10.1038/s41562-025-02267-6. URL <http://dx.doi.org/10.1038/s41562-025-02267-6>.
- [110] Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. tinybenchmarks: evaluating llms with fewer examples. In *Proceedings of the 41st International Conference on Machine Learning*, ICML'24. JMLR.org, 2024.
 - [111] Vinodkumar Prabhakaran, Christopher Homan, Lora Aroyo, Aida Mostafazadeh Davani, Alicia Parrish, Alex Taylor, Mark Diaz, Ding Wang, and Gregory Serapio-García. GRASP: A disagreement analysis framework to assess group associations in perspectives. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3473–3492, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.naacl-long.190. URL <https://aclanthology.org/2024.naacl-long.190/>.
 - [112] Tianyi Qiu, Yang Zhang, Xuchuan Huang, Jasmine Xinze Li, Jiaming Ji, and Yaodong Yang. Progressgym: Alignment with a millennium of moral progress. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 14570–14607. Curran Associates, Inc., 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/1a6d49c1a298ebb799d005b7b90ab31d-Paper-Datasets_and_Benchmarks_Track.pdf.
 - [113] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 8748–8763. PMLR, 2021. URL <https://proceedings.mlr.press/v139/radford21a.html>.
 - [114] Ahmad Radwan, Mohannad Amarneh, Hussam Alawneh, Huthaifa I. Ashqar, Anas AlSobeh, and Aws Abed Al Raheem Magableh. Predictive analytics in mental health leveraging llm embeddings and machine learning models for social media analysis. *International Journal of Web Services Research*, 21(1):1–22, February 2024. ISSN 1546-5004. doi: 10.4018/ijwsr.338222. URL <http://dx.doi.org/10.4018/IJWSR.338222>.
 - [115] Georg Rasch. *Probabilistic models for some intelligence and attainment tests*. ERIC, 1993. URL <https://eric.ed.gov/?id=ED419814>.
 - [116] Pooya Razavi and Sonya J. Powers. Estimating item difficulty using large language models and tree-based machine learning algorithms, 2025. URL <https://arxiv.org/abs/2504.08804>.
 - [117] Yuanyi Ren, Haoran Ye, Hanjun Fang, Xin Zhang, and Guojie Song. ValueBench: Towards comprehensively evaluating value orientations and understanding of large language models. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2015–2040, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.111. URL <https://aclanthology.org/2024.acl-long.111/>.
 - [118] Paul Röttger, Valentin Hofmann, Valentina Pyatkin, Musashi Hinck, Hannah Kirk, Hinrich Schuetze, and Dirk Hovy. Political compass or spinning arrow? towards more meaningful evaluations for values and opinions in large language models. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15295–15311, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.816. URL <https://aclanthology.org/2024.acl-long.816/>.
 - [119] L. L. Russell-Lasalandra, Alexander P. Christensen, and Hudson Golino. Generative psychometrics via ai-genie: Automatic item generation and validation via network-integrated evaluation, 2024. URL <https://doi.org/10.31234/osf.io/fgbj4>.
 - [120] Aadesh Salecha, Molly E Ireland, Shashanka Subrahmanya, João Sedoc, Lyle H Ungar, and Johannes C Eichstaedt. Large language models display human-like social desirability biases

- in big five personality surveys. *PNAS Nexus*, 3(12), November 2024. ISSN 2752-6542. doi: 10.1093/pnasnexus/pgae533. URL <http://dx.doi.org/10.1093/pnasnexus/pgae533>.
- [121] Alexander Scarlatos, Nigel Fernandez, Christopher Ormerod, Susan Lottridge, and Andrew Lan. Smart: Simulated students aligned with item response theory for question difficulty prediction, 2025. URL <https://arxiv.org/abs/2507.05129>.
 - [122] Greg Serapio-García, Mustafa Safdari, Clément Crepy, Luning Sun, Stephen Fitz, Peter Romero, Marwa Abdulhai, Aleksandra Faust, and Maja Matarić. Personality traits in large language models, 2023. URL <https://arxiv.org/abs/2307.00184>.
 - [123] Daun Shin, Hyoseung Kim, Seunghwan Lee, Younhee Cho, and Whanbo Jung. Using large language models to detect depression from user-generated diary text data as a novel approach in digital mental health screening: Instrument validation study. *Journal of Medical Internet Research*, 26:e54617, September 2024. ISSN 1438-8871. doi: 10.2196/54617. URL <http://dx.doi.org/10.2196/54617>.
 - [124] Oleh Shliazhko, Alena Fenogenova, Maria Tikhonova, Anastasia Kozlova, Vladislav Mikhailov, and Tatiana Shavrina. mgpt: Few-shot learners go multilingual. *Transactions of the Association for Computational Linguistics*, 12:58–79, 01 2024. ISSN 2307-387X. doi: 10.1162/tac1_a_00633. URL https://doi.org/10.1162/tac1_a_00633.
 - [125] Bangzhao Shu, Lechen Zhang, Minje Choi, Lavinia Dunagan, Lajanugen Logeswaran, Moon-tae Lee, Dallas Card, and David Jurgens. You don’t need a personality test to know these models are unreliable: Assessing the reliability of large language models on psychometric instruments. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, page 5263–5281. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.naacl-long.295. URL <http://dx.doi.org/10.18653/v1/2024.naacl-long.295>.
 - [126] Gabriel Simmons and Christopher Hare. Large language models as subpopulation representative models: A review, 2023. URL <https://arxiv.org/abs/2310.17888>.
 - [127] Joseph J. P. Simons, Wong Liang Ze, Prasanta Bhattacharya, Brandon Siyuan Loh, and Wei Gao. From traces to measures: Large language models as a tool for psychological measurement from text, 2024. URL <https://arxiv.org/abs/2405.07447>.
 - [128] Pushpdeep Singh, Mayur Patidar, and Lovekesh Vig. Translating across cultures: Llms for intralingual cultural adaptation. In *Proceedings of the 28th Conference on Computational Natural Language Learning*, page 400–418. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.conll-1.30. URL <http://dx.doi.org/10.18653/v1/2024.conll-1.30>.
 - [129] Anand Siththaranjan, Cassidy Laidlaw, and Dylan Hadfield-Menell. Distributional preference learning: Understanding and accounting for hidden context in rlhf, 2024. URL <https://arxiv.org/abs/2312.08358>.
 - [130] Sarath Sivaprasad, Pramod Kaushik, Sahar Abdelnabi, and Mario Fritz. A theory of response sampling in LLMs: Part descriptive and part prescriptive. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 30091–30135, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. doi: 10.18653/v1/2025.acl-long.1454. URL <https://aclanthology.org/2025.acl-long.1454/>.
 - [131] Taylor Sorensen, Pushkar Mishra, Roma Patel, Michael Henry Tessler, Michiel Bakker, Georgina Evans, Iason Gabriel, Noah Goodman, and Verena Rieser. Value profiles for encoding human variation, 2025. URL <https://arxiv.org/abs/2503.15484>.
 - [132] Joseph Suh, Suhong Moon, Minwoo Kang, and David Chan. Rediscovering the latent dimensions of personality with large language models as trait descriptors. In *NeurIPS 2024 Workshop on Behavioral Machine Learning*, 2024. URL <https://openreview.net/forum?id=HmtCakUGHj>.
 - [133] Youmi Suk and Kyung T. Han. A psychometric framework for evaluating fairness in algorithmic decision making: Differential algorithmic functioning. *Journal of Educational and Behavioral Statistics*, 49(2):151–172, May 2023. ISSN 1935-1054. doi: 10.3102/10769986231171711. URL <http://dx.doi.org/10.3102/10769986231171711>.

- [134] Kyle Swanson, Wesley Wu, Nash L. Bulaong, John E. Pak, and James Zou. The virtual lab of ai agents designs new sars-cov-2 nanobodies. *Nature*, July 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-09442-9. URL <http://dx.doi.org/10.1038/s41586-025-09442-9>.
- [135] Andreas Säuberli, Diego Frassinelli, and Barbara Plank. Do llms give psychometrically plausible responses in educational assessments?, 2025. URL <https://arxiv.org/abs/2506.09796>.
- [136] Tom Sühr, Florian E. Dorner, Samira Samadi, and Augustin Kelava. Challenging the validity of personality tests for large language models, 2024. URL <https://arxiv.org/abs/2311.05297>.
- [137] Alex Tamkin, Amanda Askell, Liane Lovitt, Esin Durmus, Nicholas Joseph, Shauna Kravec, Karina Nguyen, Jared Kaplan, and Deep Ganguli. Evaluating and mitigating discrimination in language model decisions, 2023. URL <https://arxiv.org/abs/2312.03689>.
- [138] Gemini Team, Rohan Anil, Sebastian Borgeaud, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, Katie Millican, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023.
- [139] Yuto Tomikawa, Ayaka Suzuki, and Masaki Uto. Adaptive question–answer generation with difficulty control using item response theory and pretrained transformer models. *IEEE Transactions on Learning Technologies*, 17:2186–2198, 2024. ISSN 2372-0050. doi: 10.1109/tlt.2024.3491801. URL <http://dx.doi.org/10.1109/TLT.2024.3491801>.
- [140] Song Tong, Kai Mao, Zhen Huang, Yukun Zhao, and Kaiping Peng. Automating psychological hypothesis generation with ai: when large language models meet causal graph. *Humanities and Social Sciences Communications*, 11(1), July 2024. ISSN 2662-9992. doi: 10.1057/s41599-024-03407-5. URL <http://dx.doi.org/10.1057/s41599-024-03407-5>.
- [141] Sang Truong, Yuheng Tu, Percy Liang, Bo Li, and Sanmi Koyejo. Reliable and efficient amortized model-based evaluation, 2025. URL <https://arxiv.org/abs/2503.13335>.
- [142] Sang Truong, Yuheng Tu, Percy Liang, Bo Li, and Sanmi Koyejo. Reliable and efficient amortized model-based evaluation, 2025. URL <https://arxiv.org/abs/2503.13335>.
- [143] Adithya V Ganesan, Yash Kumar Lal, August Nilsson, and H. Andrew Schwartz. Systematic evaluation of GPT-3 for zero-shot personality estimation. In Jeremy Barnes, Orphée De Clercq, and Roman Klinger, editors, *Proceedings of the 13th Workshop on Computational Approaches to Subjectivity, Sentiment, & Social Media Analysis*, pages 390–400, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.wassa-1.34. URL <https://aclanthology.org/2023.wassa-1.34/>.
- [144] Danny Valdez, Andrew C. Pickett, and Patricia Goodson. Topic modeling: Latent semantic analysis for the social sciences. *Social Science Quarterly*, 99(5):1665–1679, September 2018. ISSN 1540-6237. doi: 10.1111/ssqu.12528. URL <http://dx.doi.org/10.1111/ssqu.12528>.
- [145] Vasudha Varadarajan, Sverker Sikström, Oscar Kjell, and H. Andrew Schwartz. ALBA: Adaptive language-based assessments for mental health. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 2466–2478, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.naacl-long.136. URL <https://aclanthology.org/2024.naacl-long.136/>.
- [146] Alina A von Davier, Robert J Mislevy, and Jiangang Hao. *Computational psychometrics: New methodologies for a new generation of digital learning and assessment: With examples in R and Python*. Springer Nature, 2022. URL <https://link.springer.com/book/10.1007/978-3-030-74394-9>.
- [147] Angelina Wang, Jamie Morgenstern, and John P. Dickerson. Large language models that replace human participants can harmfully misportray and flatten identity groups. *Nature Machine Intelligence*, 7(3):400–411, February 2025. ISSN 2522-5839. doi: 10.1038/s42256-025-00986-z. URL <http://dx.doi.org/10.1038/s42256-025-00986-z>.
- [148] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, brian ichter, Fei Xia, Ed Chi, Quoc V Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language

- models. In S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837. Curran Associates, Inc., 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/file/9d5609613524ecf4f15af0f7b31abca4-Paper-Conference.pdf.
- [149] Aaron C. Weidman, Conor M. Steckler, and Jessica L. Tracy. The jingle and jangle of emotion assessment: Imprecise measurement, casual scale usage, and conceptual fuzziness in emotion research. *Emotion*, 17(2):267–295, 2017. ISSN 1528-3542. doi: 10.1037/emo0000226. URL <http://dx.doi.org/10.1037/emo0000226>.
 - [150] Xuansheng Wu, Padmaja Pravin Saraf, Gyeonggeon Lee, Ehsan Latif, Ninghao Liu, and Xiaoming Zhai. Unveiling scoring processes: Dissecting the differences between llms and human graders in automatic scoring, 2025. URL <https://arxiv.org/abs/2407.18328>.
 - [151] Dirk U. Wulff and Rui Mata. Semantic embeddings reveal and address taxonomic incommensurability in psychological measurement. *Nature Human Behaviour*, 9(5):944–954, March 2025. ISSN 2397-3374. doi: 10.1038/s41562-024-02089-y. URL <http://dx.doi.org/10.1038/s41562-024-02089-y>.
 - [152] Xuhai Xu, Bingsheng Yao, Yuanzhe Dong, Saadia Gabriel, Hong Yu, James Hendler, Marzyeh Ghassemi, Anind K. Dey, and Dakuo Wang. Mental-llm: Leveraging large language models for mental health prediction via online text data. *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.*, 8(1), 2024. doi: 10.1145/3643540. URL <https://doi.org/10.1145/3643540>.
 - [153] Zhenrong Xu, Yuan Gao, Fang Wang, Longqian Zhang, Li Zhang, Junke Wang, and Jie Shu. Depression detection methods based on multimodal fusion of voice and text. *Scientific Reports*, 15(1), July 2025. ISSN 2045-2322. doi: 10.1038/s41598-025-03524-4. URL <http://dx.doi.org/10.1038/s41598-025-03524-4>.
 - [154] Yang Yan, Lizhi Ma, Anqi Li, Jingsong Ma, and Zhenzhong Lan. Predicting the big five personality traits in chinese counselling dialogues using large language models, 2024. URL <https://arxiv.org/abs/2406.17287>.
 - [155] Qisen Yang, Zekun Wang, Honghui Chen, Shenzi Wang, Yifan Pu, Xin Gao, Wenhao Huang, Shiji Song, and Gao Huang. PsychoGAT: A novel psychological measurement paradigm through interactive fiction games with LLM agents. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14470–14505, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.779. URL <https://aclanthology.org/2024.acl-long.779/>.
 - [156] Tao Yang, Tianyuan Shi, Fanqi Wan, Xiaojun Quan, Qifan Wang, Bingzhe Wu, and Jiaxiang Wu. PsyCoT: Psychological questionnaire as powerful chain-of-thought for personality detection. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 3305–3320, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.216. URL <https://aclanthology.org/2023.findings-emnlp.216/>.
 - [157] Haoran Ye, Jing Jin, Yuhang Xie, Xin Zhang, and Guojie Song. Large language model psychometrics: A systematic review of evaluation, validation, and enhancement, 2025. URL <https://arxiv.org/abs/2505.08245>.
 - [158] Haoran Ye, Yuhang Xie, Yuanyi Ren, Hanjun Fang, Xin Zhang, and Guojie Song. Measuring human and ai values based on generative psychometrics with large language models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(25):26400–26408, 2025. doi: 10.1609/aaai.v39i25.34839. URL <https://ojs.aaai.org/index.php/AAAI/article/view/34839>.
 - [159] Haoran Ye, Tianze Zhang, Yuhang Xie, Liyuan Zhang, Yuanyi Ren, Xin Zhang, and Guojie Song. Generative psycho-lexical approach for constructing value systems in large language models, 2025. URL <https://arxiv.org/abs/2502.02444>.
 - [160] Junhao Yu, Yan Zhuang, YuXuan Sun, Weibo Gao, Qi Liu, Mingyue Cheng, Zhenya Huang, and Enhong Chen. Testagent: An adaptive and intelligent expert for human assessment, 2025. URL <https://arxiv.org/abs/2506.03032>.
 - [161] Mert Yuksekgonul, Federico Bianchi, Joseph Boen, Sheng Liu, Pan Lu, Zhi Huang, Carlos Guestrin, and James Zou. Optimizing generative ai by backpropagating language model

- feedback. *Nature*, 639(8055):609–616, March 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-08661-4. URL <http://dx.doi.org/10.1038/s41586-025-08661-4>.
- [162] Eric Zelikman, Wanjing Ma, Jasmine Tran, Diyi Yang, Jason Yeatman, and Nick Haber. Generating and evaluating tests for k-12 students with language model simulations: A case study on sentence reading efficiency. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2190–2205, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.135. URL <https://aclanthology.org/2023.emnlp-main.135/>.
 - [163] Jingshen Zhang, Jiajun Xie, and Xinying Qiu. Controlling cloze-test question item difficulty with plm-based surrogate models for irt assessment, 2024. URL <https://arxiv.org/abs/2403.01456>.
 - [164] Tianyi Zhang, Antonis Koutsoumpis, Janneke K. Oostrom, Djurre Holtrop, Sina Ghassemi, and Reinout E. de Vries. Can large language models assess personality from asynchronous video interviews? a comprehensive evaluation of validity, reliability, fairness, and rating patterns. *IEEE Transactions on Affective Computing*, 15(3):1769–1785, July 2024. ISSN 2371-9850. doi: 10.1109/taffc.2024.3374875. URL <http://dx.doi.org/10.1109/TAFRC.2024.3374875>.
 - [165] Zhengye Zhang, Sirui Zhao, Shifeng Liu, Shukang Yin, Xinglong Mao, Tong Xu, and Enhong Chen. Mellm: Exploring llm-powered micro-expression understanding enhanced by subtle motion perception, 2025. URL <https://arxiv.org/abs/2505.07007>.
 - [166] Xiaoming Zhao, Yuehui Liao, Zhiwei Tang, Yicheng Xu, Xin Tao, Dandan Wang, Guoyu Wang, and Hongsheng Lu. Integrating audio and visual modalities for multimodal personality trait recognition via hybrid deep learning. *Frontiers in Neuroscience*, 16, January 2023. ISSN 1662-453X. doi: 10.3389/fnins.2022.1107284. URL <http://dx.doi.org/10.3389/fnins.2022.1107284>.
 - [167] Wenjie Zheng, Qiming Xie, Zengzhi Wang, Jianfei Yu, and Rui Xia. Towards explainable multimodal depression recognition for clinical interviews, 2025. URL <https://arxiv.org/abs/2501.16106>.
 - [168] Xuhui Zhou, Hao Zhu, Leena Mathur, Ruohong Zhang, Haoifei Yu, Zhengyang Qi, Louis-Philippe Morency, Yonatan Bisk, Daniel Fried, Graham Neubig, et al. Sotopia: Interactive evaluation for social intelligence in language agents. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=mM7VurbA4r>.
 - [169] Jian-Qiao Zhu, Joshua C. Peterson, Benjamin Enke, and Thomas L. Griffiths. Capturing the complexity of human strategic decision-making with machine learning. *Nature Human Behaviour*, June 2025. ISSN 2397-3374. doi: 10.1038/s41562-025-02230-5. URL <http://dx.doi.org/10.1038/s41562-025-02230-5>.
 - [170] Jianfeng Zhu, Ruoming Jin, and Karin G. Coifman. Can llms infer personality from real world conversations?, 2025. URL <https://arxiv.org/abs/2507.14355>.
 - [171] Qingfeng Zhuge, Han Wang, and Xuyang Chen. Twinstar: A novel design for enhanced test question generation using dual-llm engine. *Applied Sciences*, 15(6):3055, March 2025. ISSN 2076-3417. doi: 10.3390/app15063055. URL <http://dx.doi.org/10.3390/app15063055>.
 - [172] Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. Can large language models transform computational social science? *Computational Linguistics*, 50(1):237–291, 2024. ISSN 1530-9312. doi: 10.1162/coli_a_00502. URL http://dx.doi.org/10.1162/coli_a_00502.